



Semiconductor
Research
Corporation



CoCoSys
CENTER FOR THE
CO-DESIGN OF COGNITIVE SYSTEMS

BERRY: Bit Error Robustness for Energy-Efficient Reinforcement Learning-Based Autonomous Systems

Zishen Wan¹, Nandhini Chandramoorthy², Karthik Swaminathan²,
Pin-Yu Chen², Vijay Janapa Reddi³, Arijit Raychowdhury¹

¹Georgia Institute of Technology ²IBM Research ³Harvard University



HARVARD
UNIVERSITY

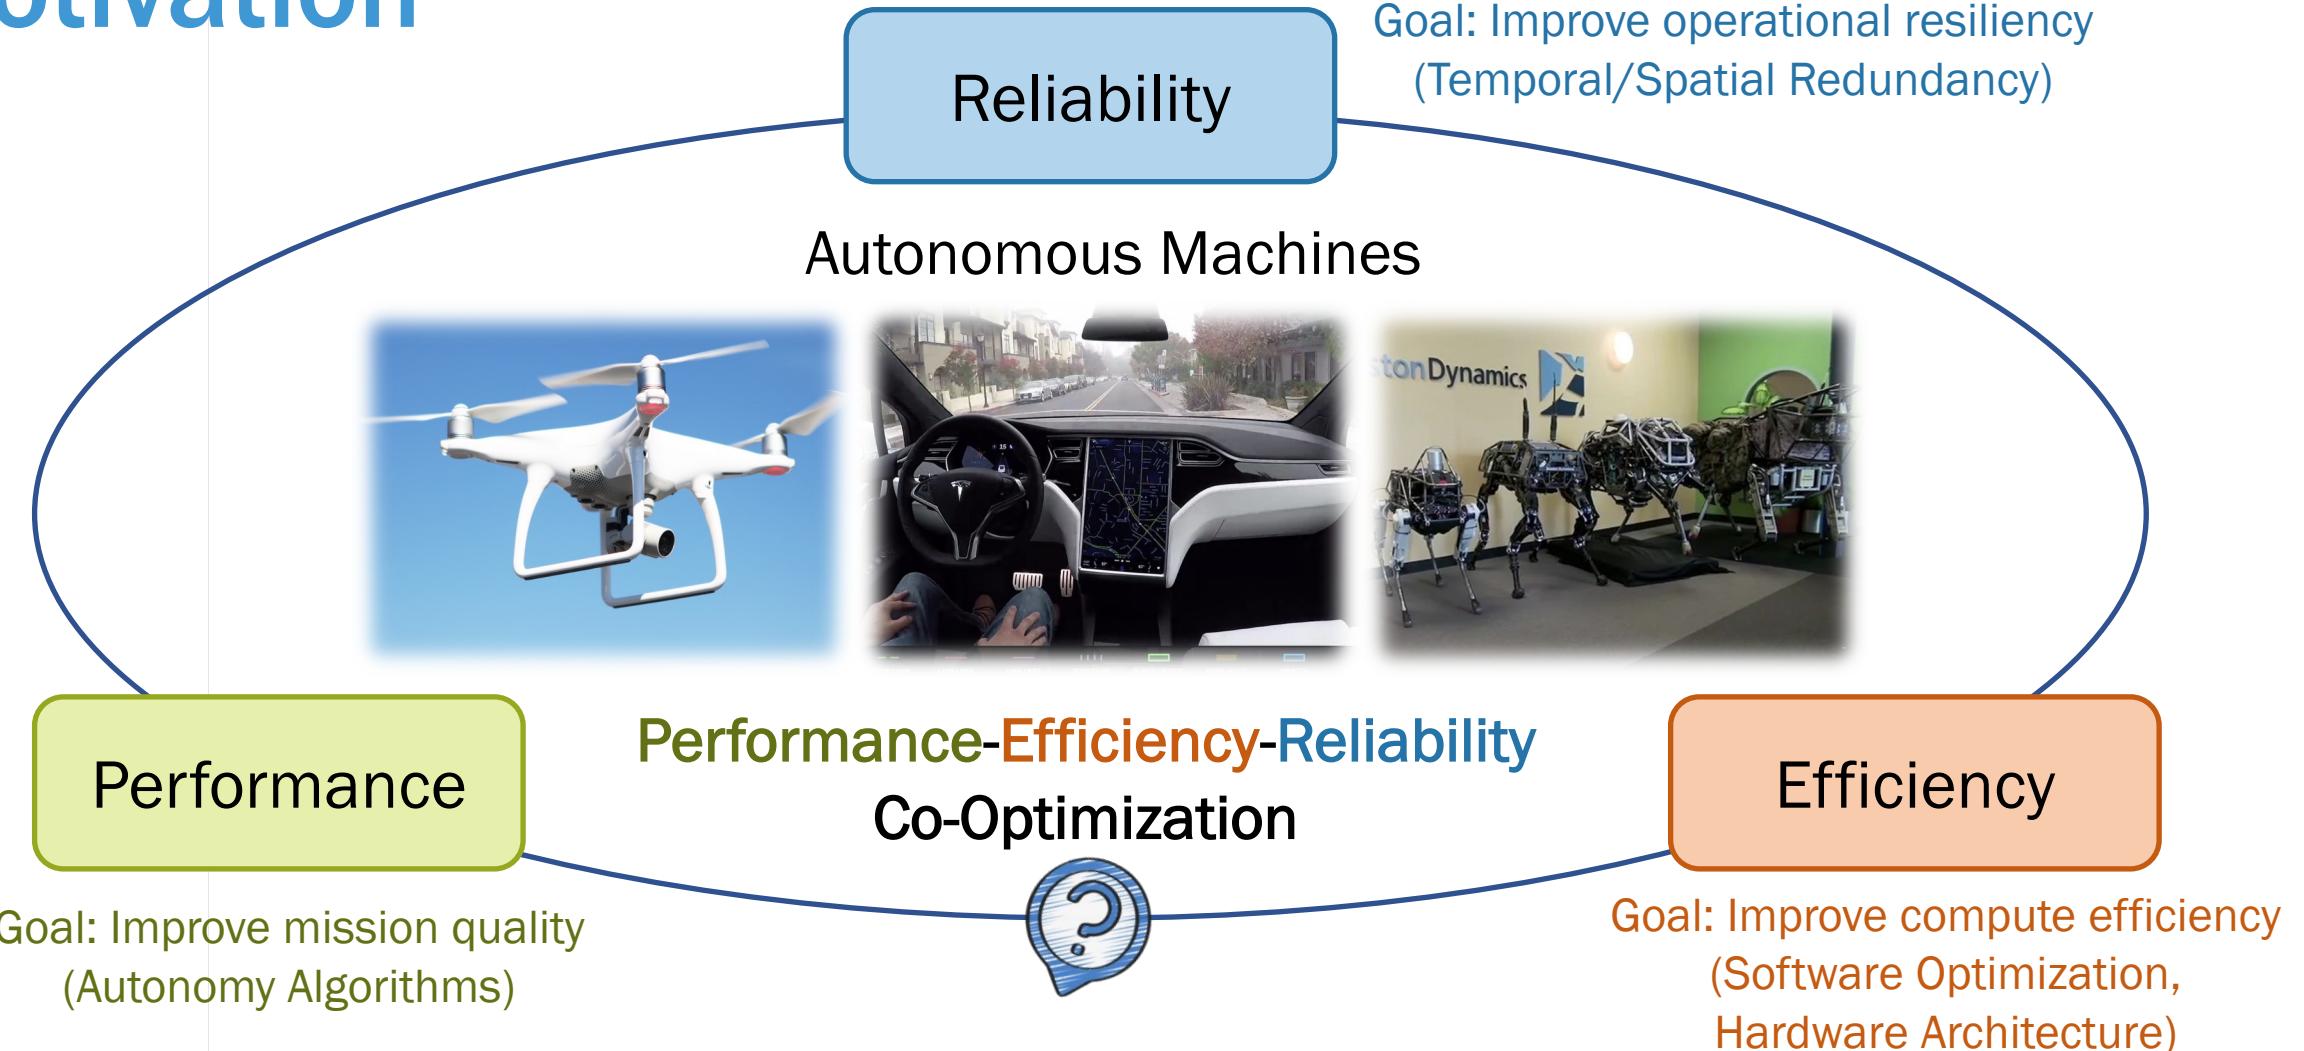


Outline

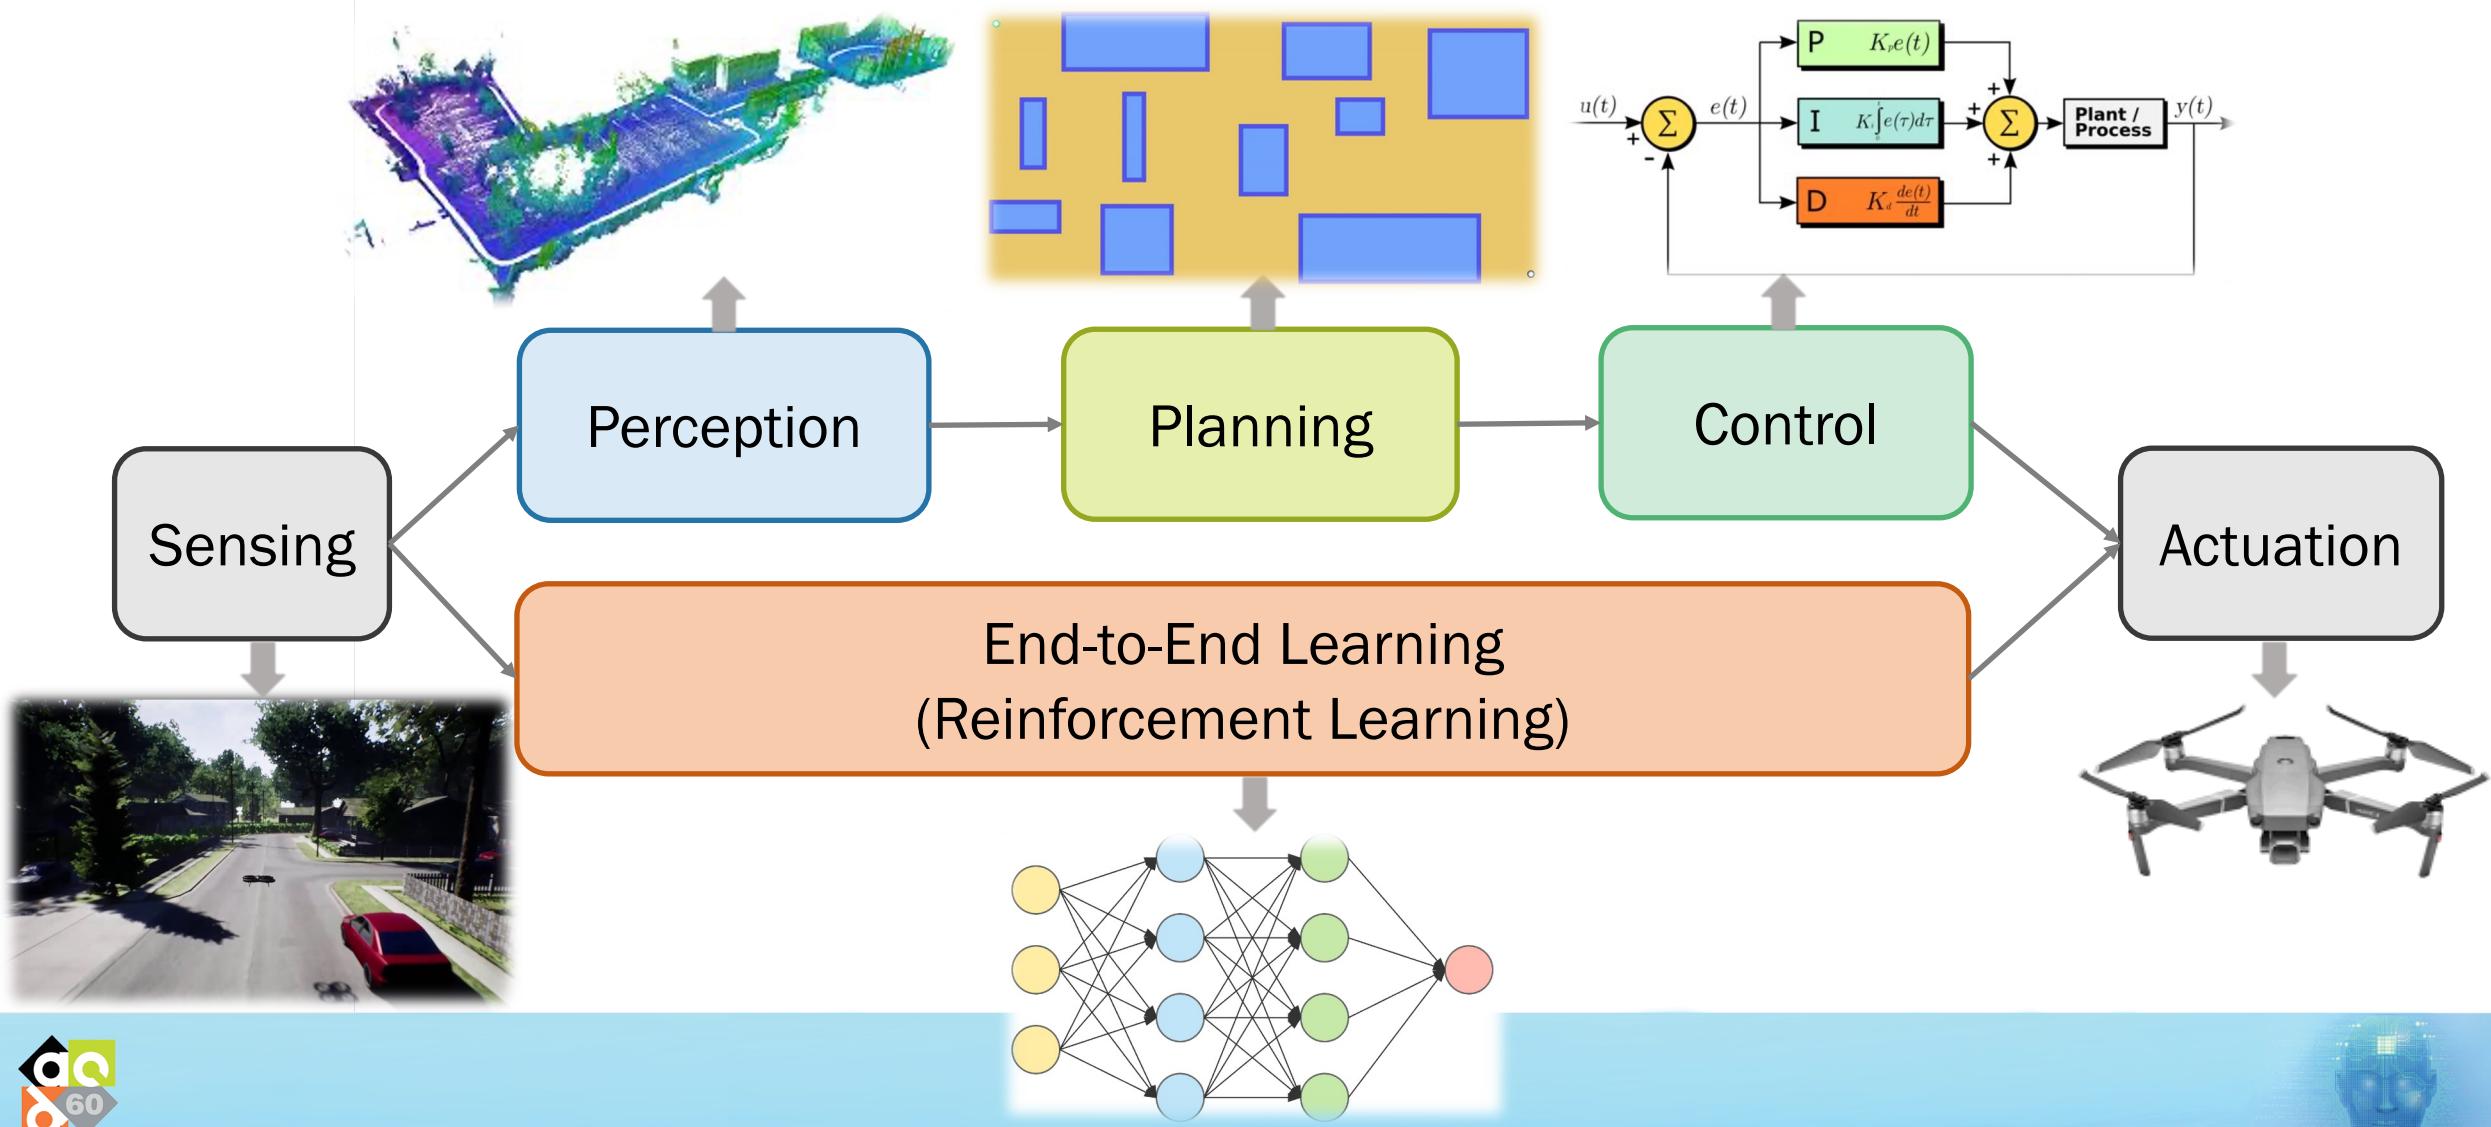
- Motivation – Why we need efficient-yet-resilient autonomous system?
- What is Autonomous Machine System
 - Physical model-based and end-to-end learning (RL)-based paradigm
- BERRY Framework
 - Drone system characterization
 - Bit-error-robustness improvement
- Evaluations
 - Various environments, models, drone platforms, chip error patterns



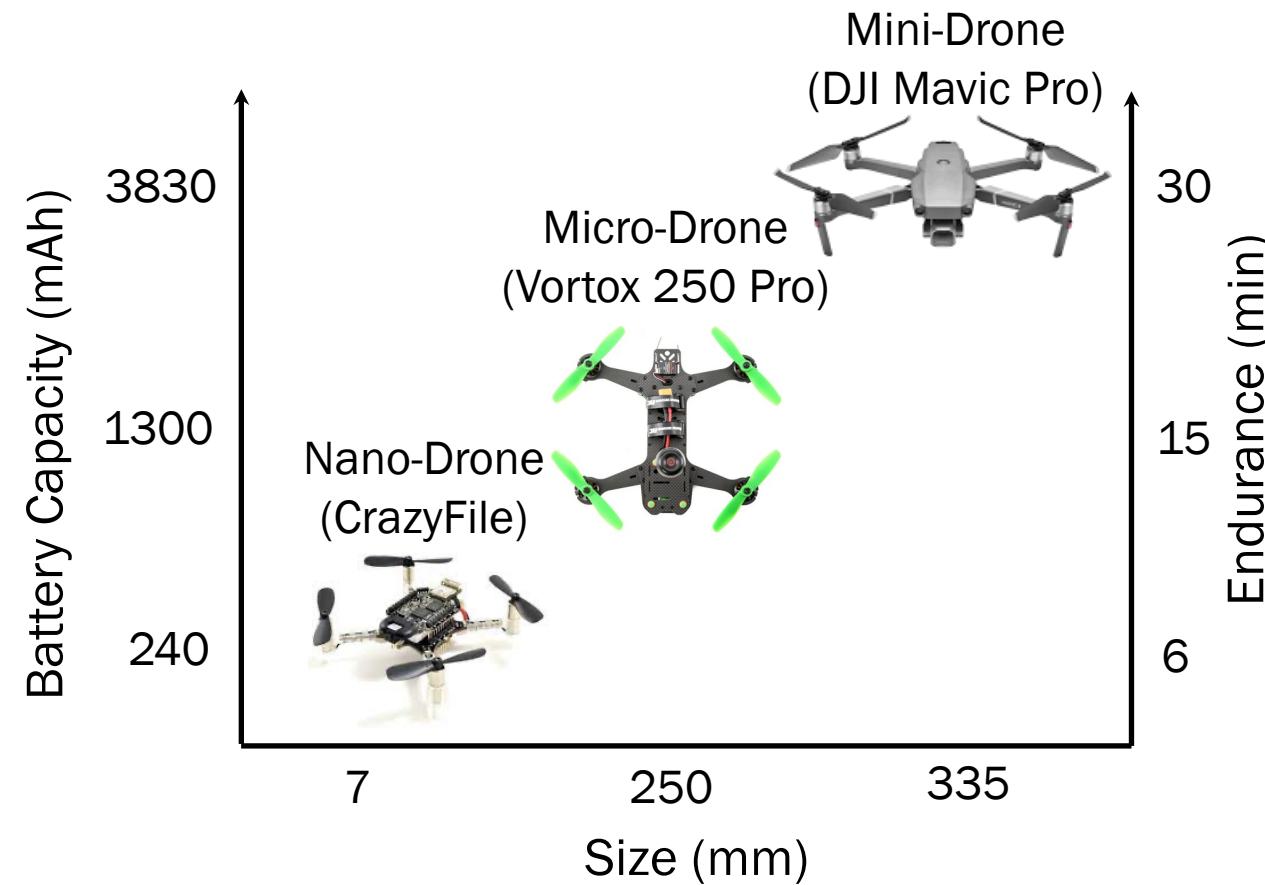
Motivation



What is Autonomous Machine System?



Challenge 1: Resource Constraints

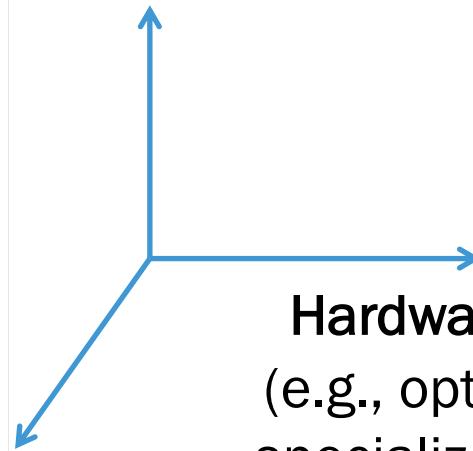


Drones are size,
weight, and power
(SWaP)
constrained



Energy-Efficient Autonomous System

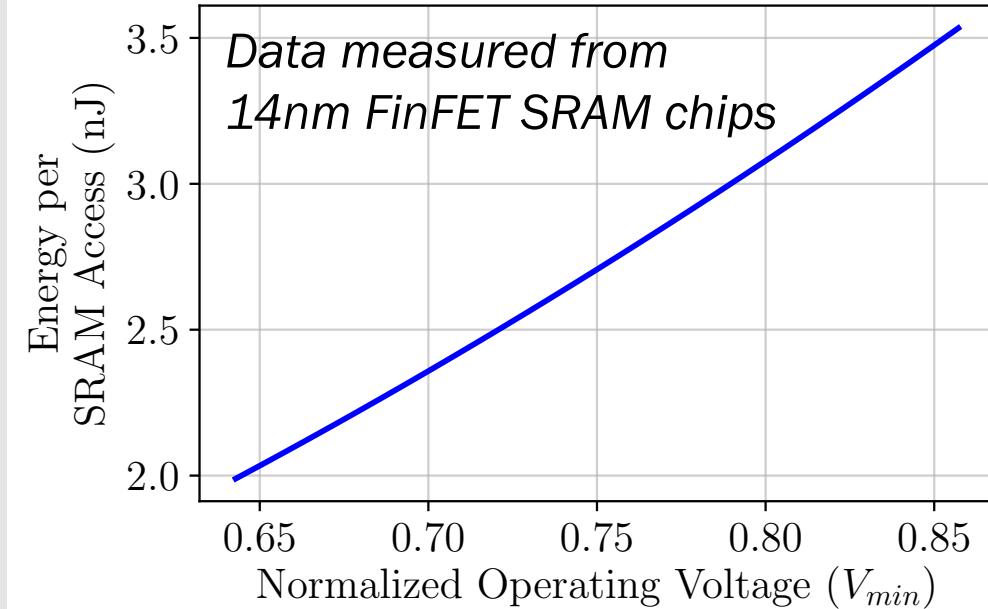
Software Optimization
(e.g., quantization, sparsity)



Lower processor
operating voltage

$$\text{Energy} \propto \text{Voltage}^2$$

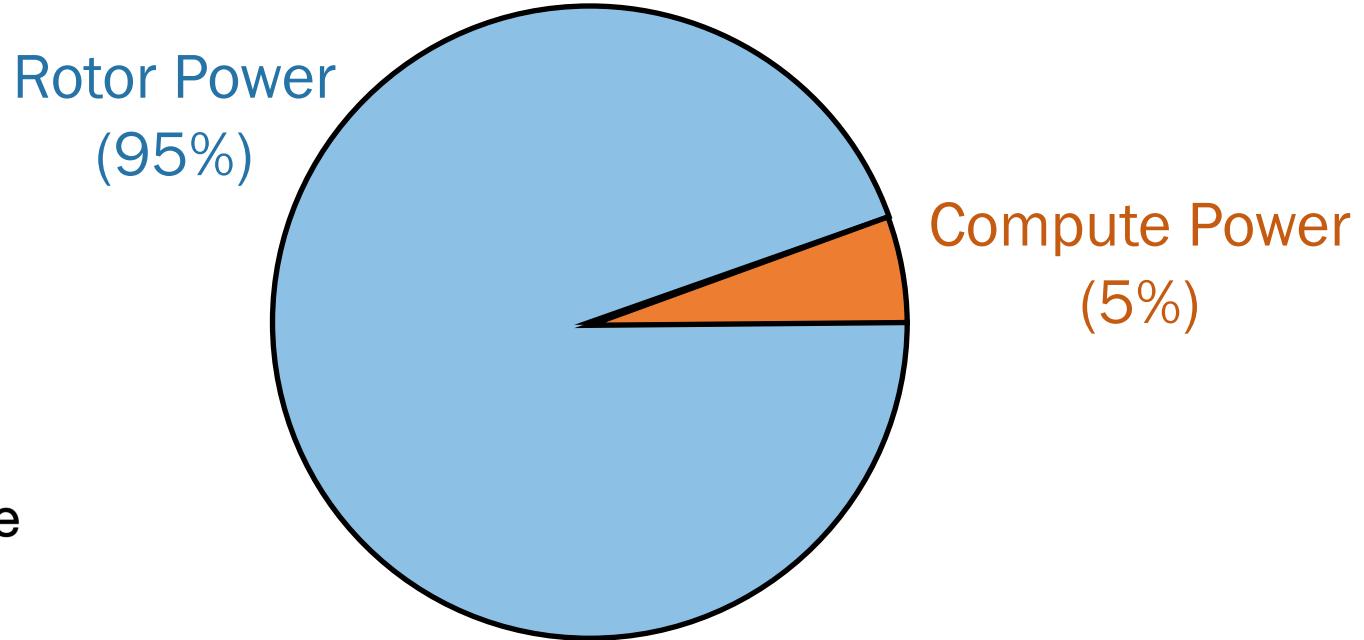
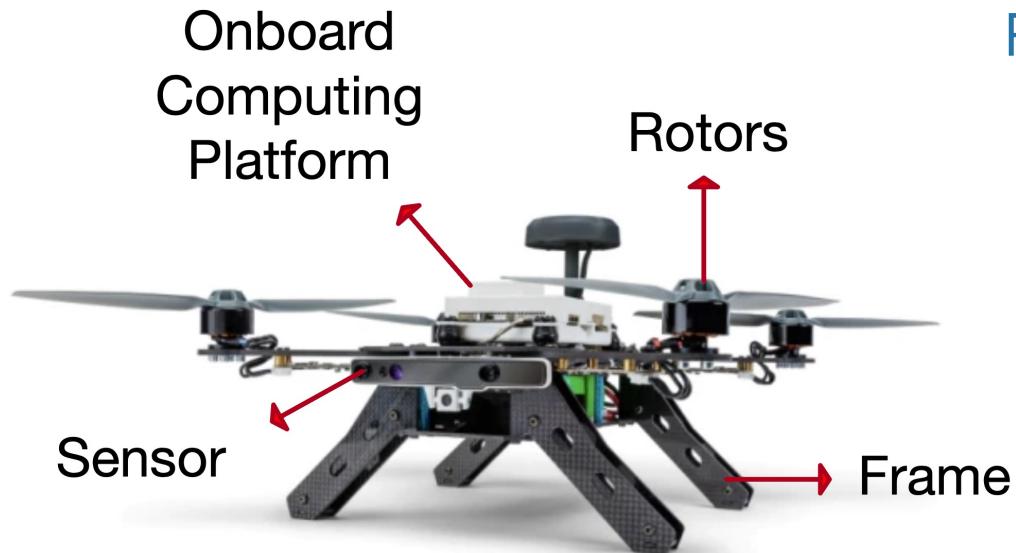
SRAM Access Energy vs. Operating Voltage



Lower operating voltage
quadratically reduces energy



Challenge 2: Compute-Physics Correlation



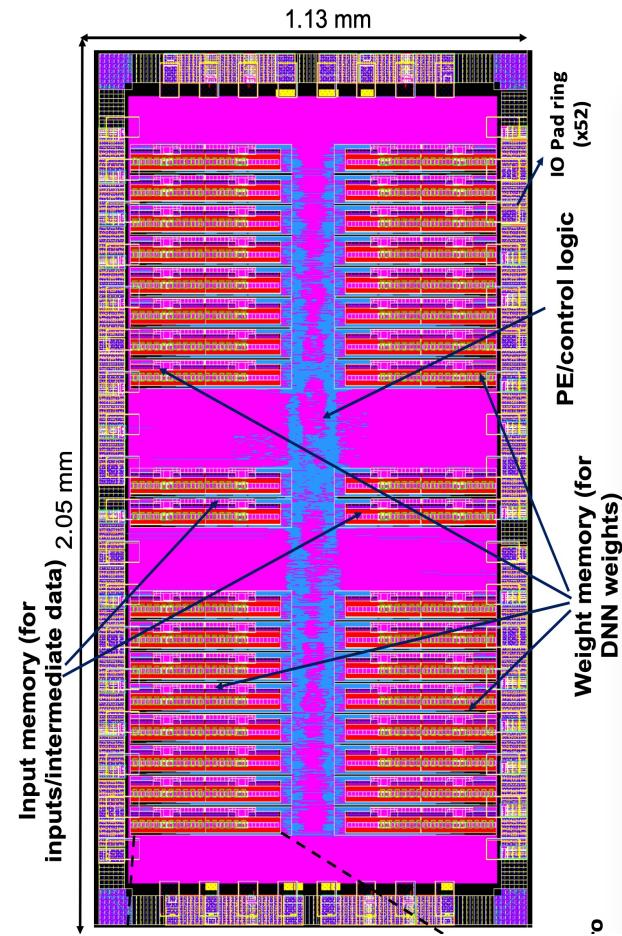
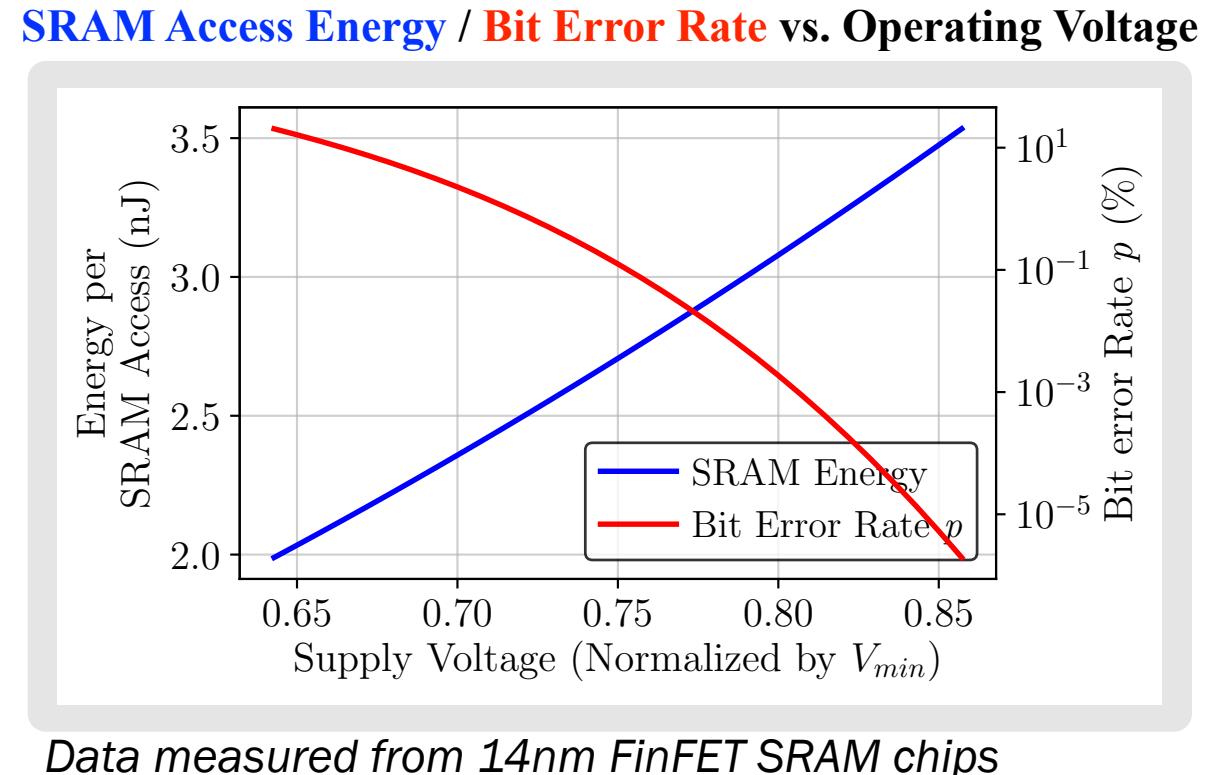
Power breakdown measured from 3DR Solo drone

Compute power is only a small fraction of total drone power
-> ***Will optimize compute bring system energy-savings?***

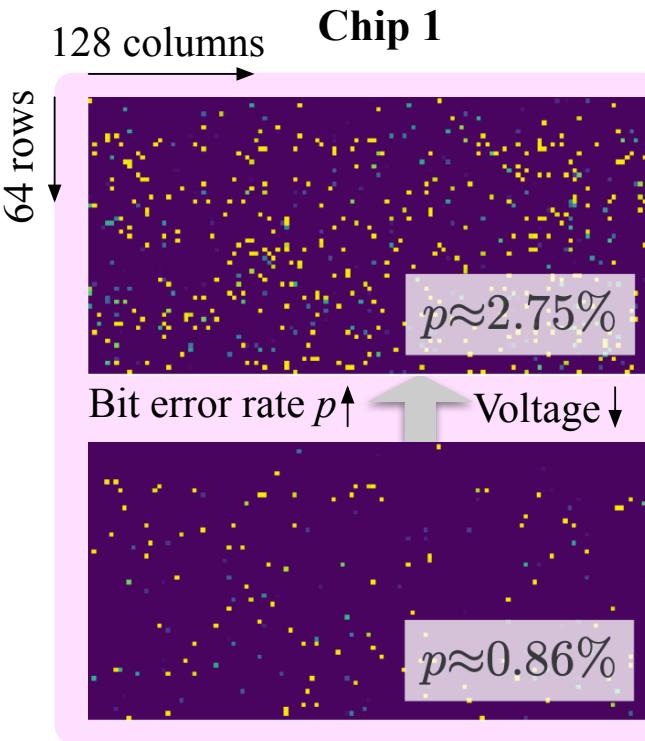


Challenge 3: Low Voltage Induces Faults

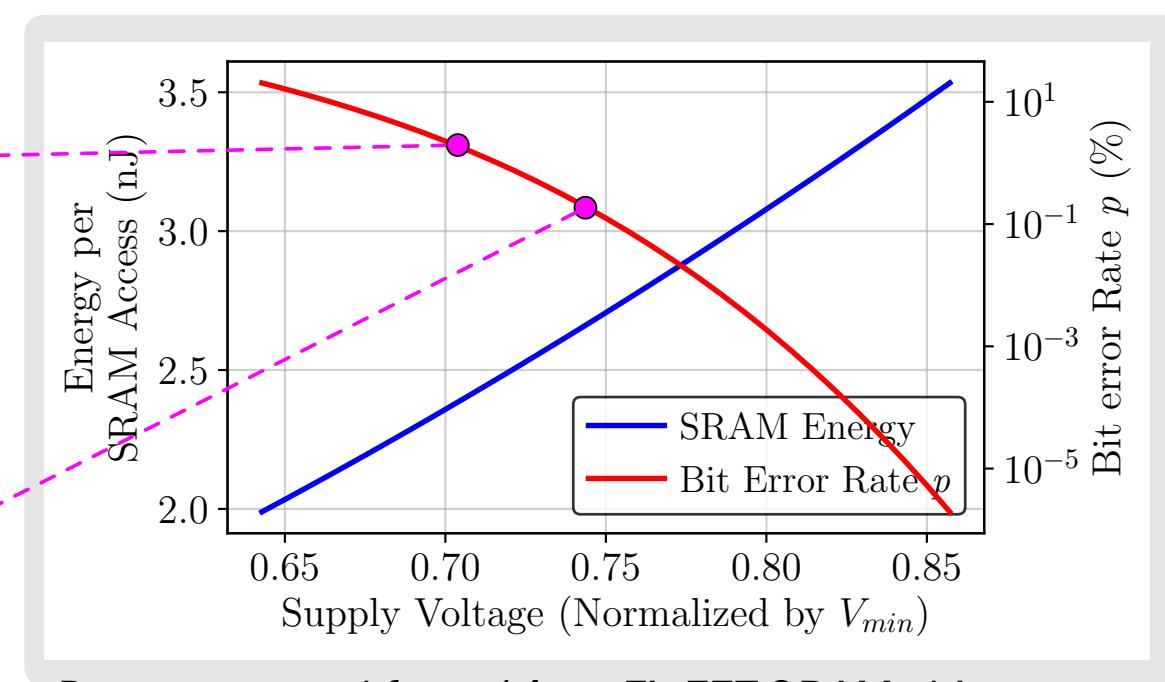
Technology	14nm
Chip Dimension	2.05 mm x 1.13 mm
Memory Capacity	128 KB weight, 16 KB input
Frequency/ Voltage	330 MHz for $V_{dd}=0.8V$



Challenge 3: Low Voltage Induces Faults



SRAM Access Energy / Bit Error Rate vs. Operating Voltage



Data measured from 14nm FinFET SRAM chips

Operating below rated voltage range results in memory bit errors, negatively impacting safety

Fault Characterization

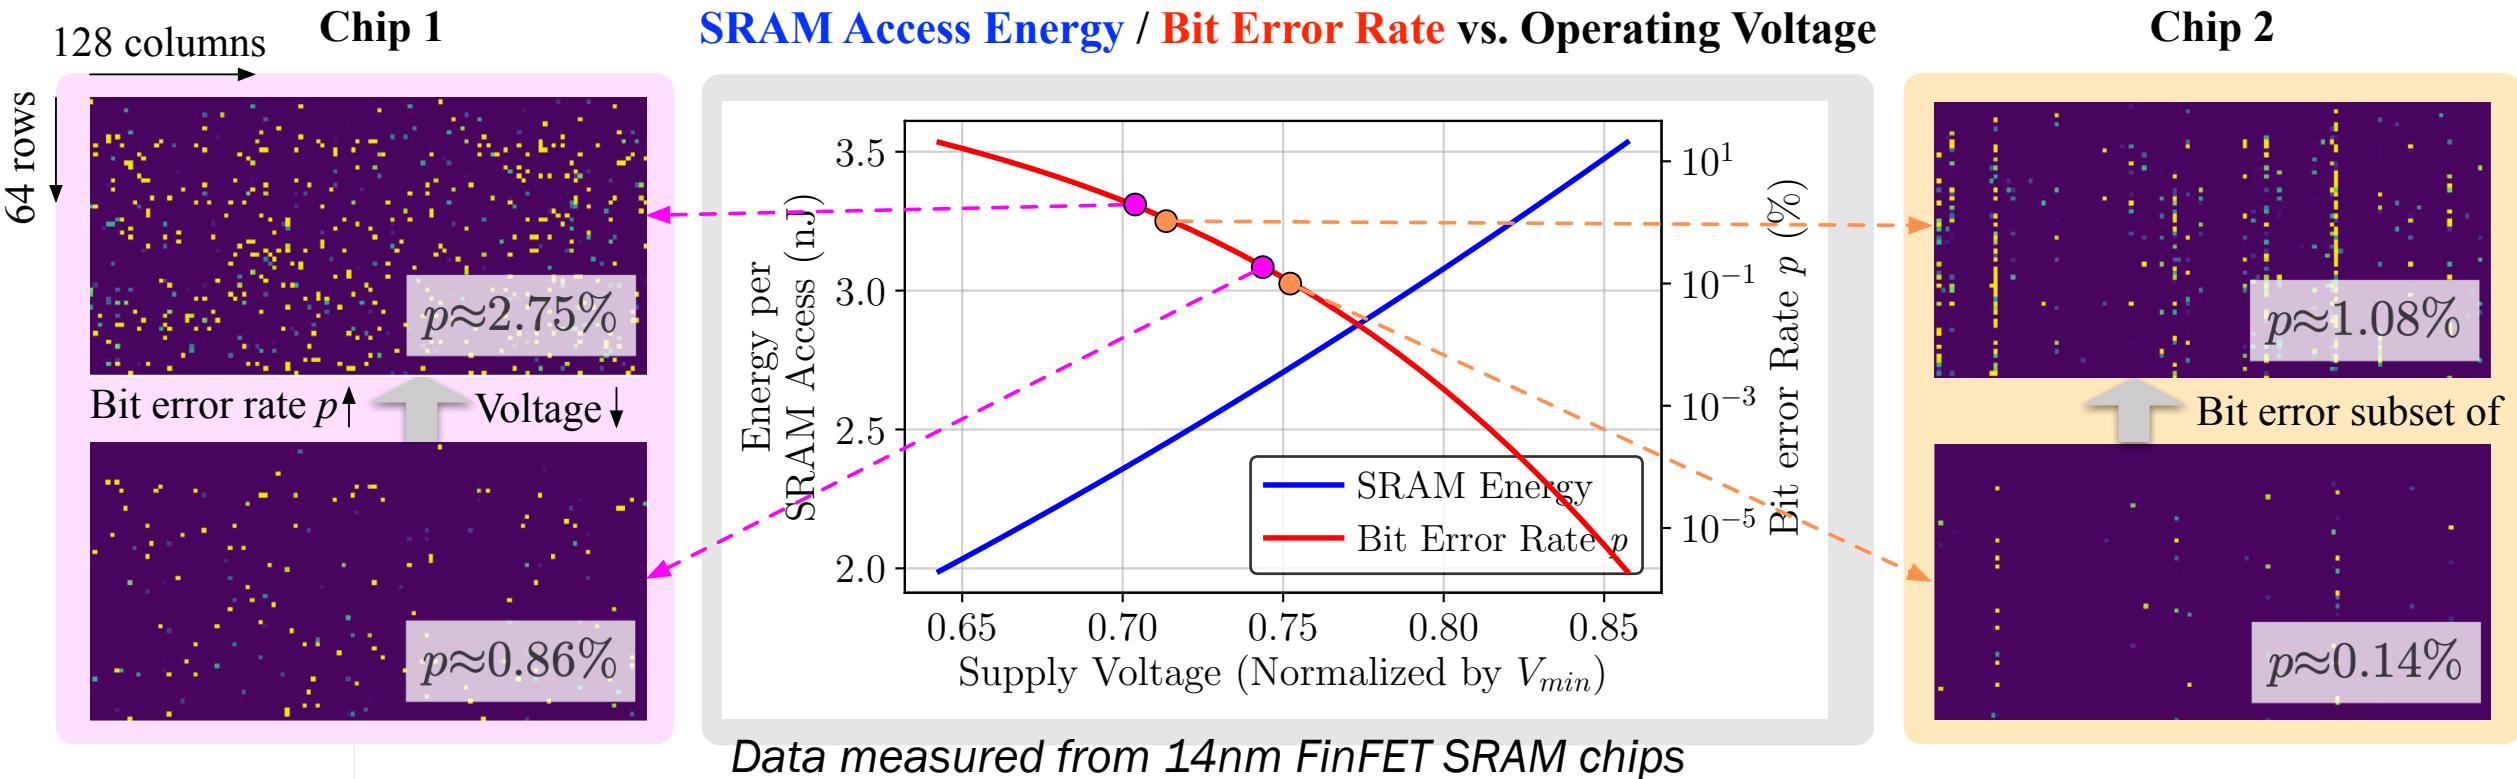
Persistent: For a fixed memory array, bit errors are persistent across voltage.

Inclusive: errors at low $BER(p)$ are subset of high $BER(p)$.

Independent: Across chips and arrays, error locations are random and independent.



Challenge 3: Low Voltage Induces Faults



Operating below rated voltage range results in memory bit errors, negatively impacting safety



BERRY

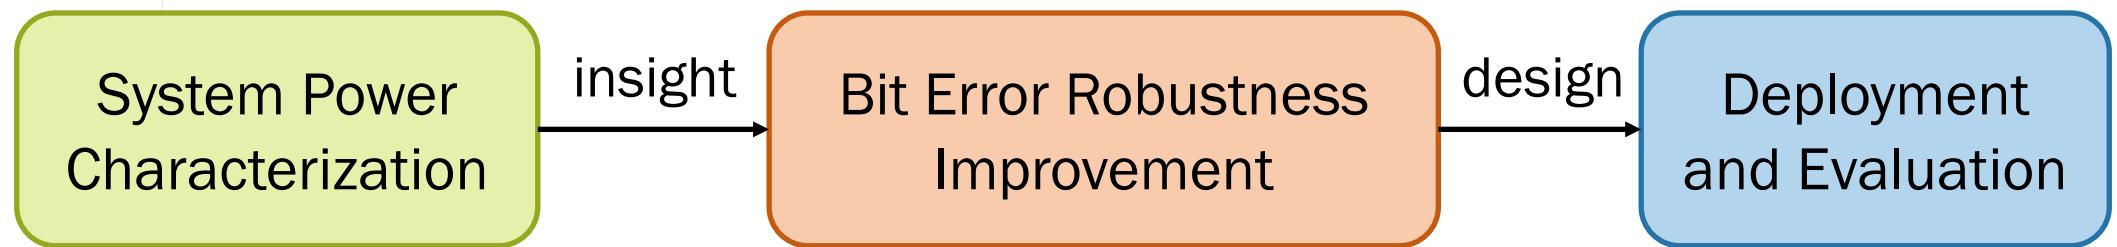
~~How to~~ enable aggressive **energy-savings** yet
computationally-resilience for autonomous system
under low-voltage operation?

(**performance-efficiency-resilience** co-optimization)



BERRY Framework

(BERRY: Bit Error Robustness for Energy-Efficient Reinforcement Learning-Based Autonomous Systems)

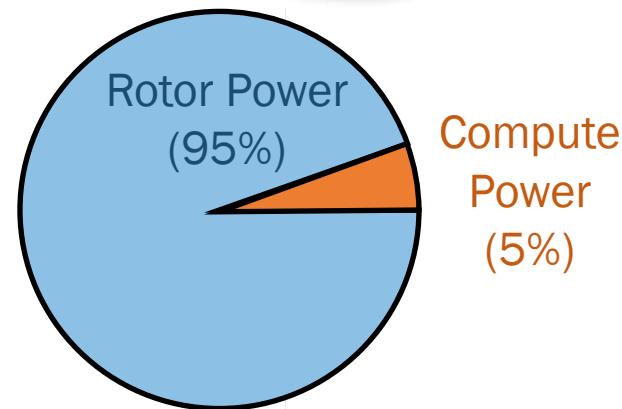
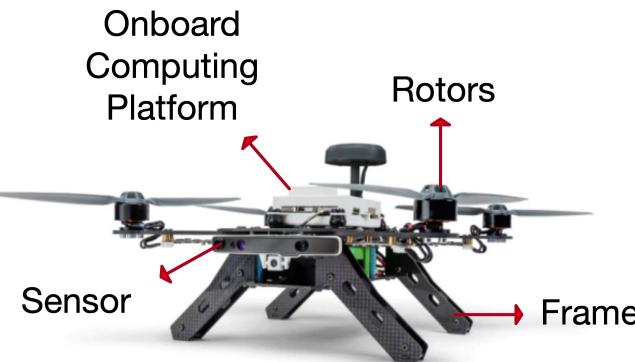


BERRY Framework

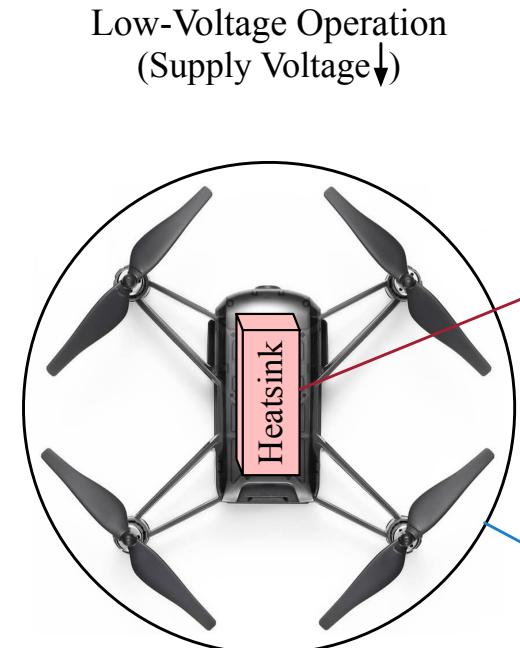
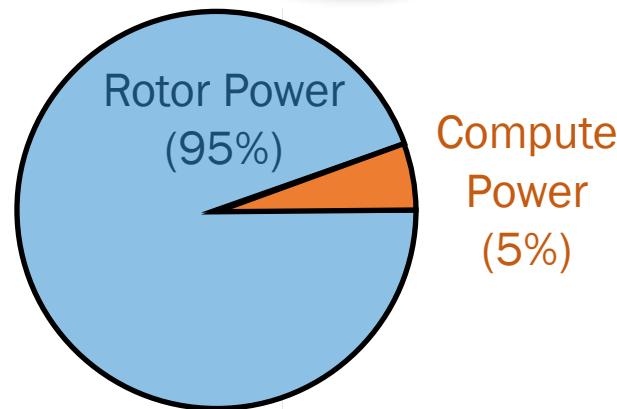
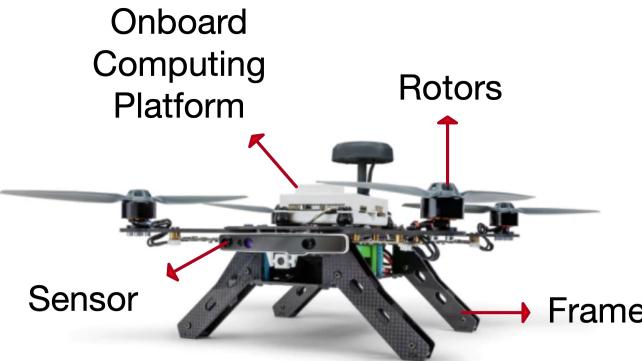
(BERRY: Bit Error Robustness for Energy-Efficient Reinforcement Learning-Based Autonomous Systems)



Small Compute Power, Large System Impact!



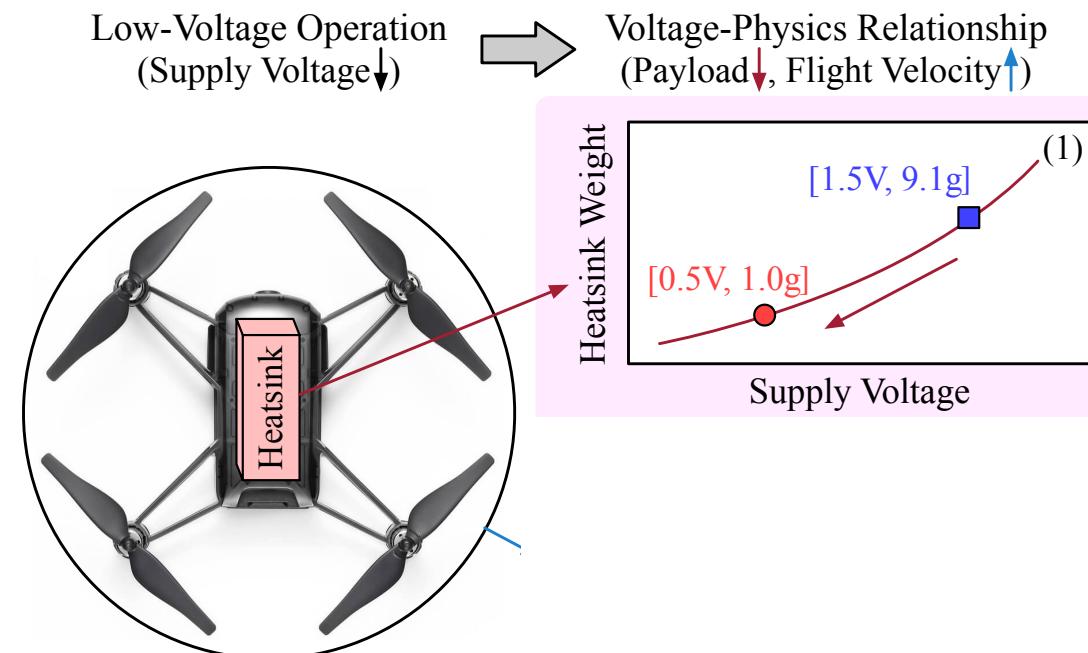
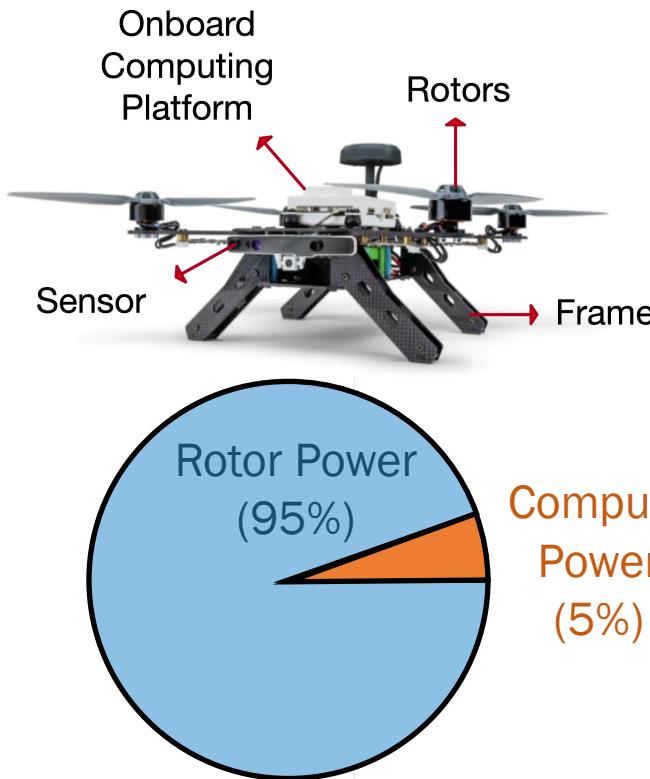
Small Compute Power, Large System Impact!



Low-voltage operation



Small Compute Power, Large System Impact!



Low-voltage operation → Peak temperature ↓, heatsink size and weight ↓
(HotSpot analysis^[1] + heatsink modeling^[2])

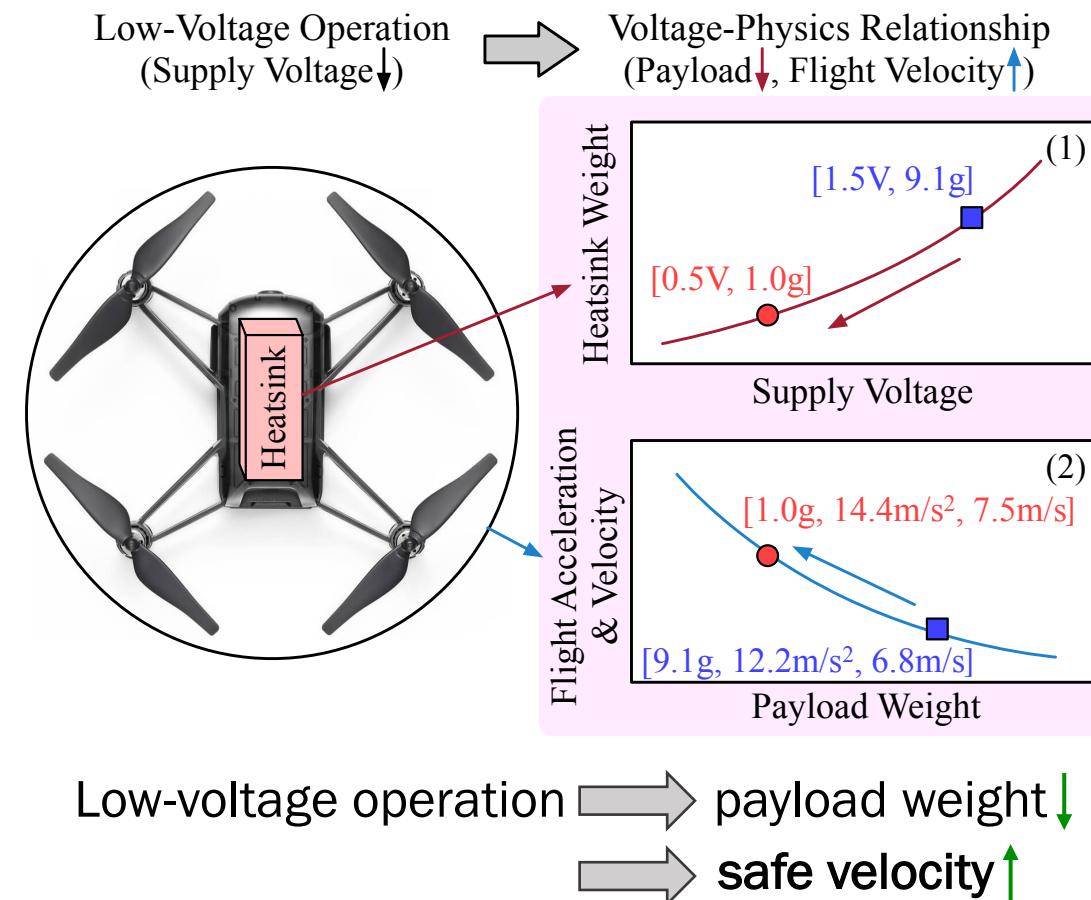
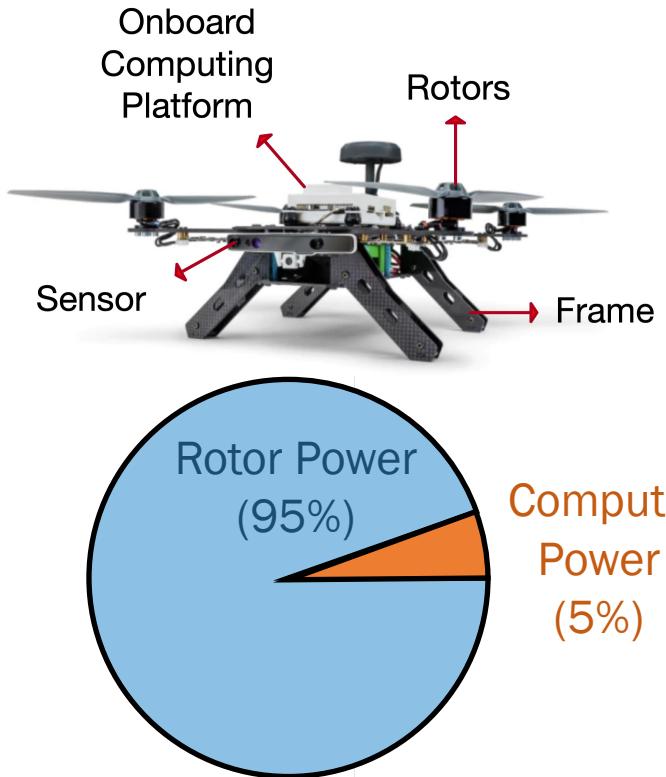


[1] Hotspot 6.0: Validation, acceleration and extension

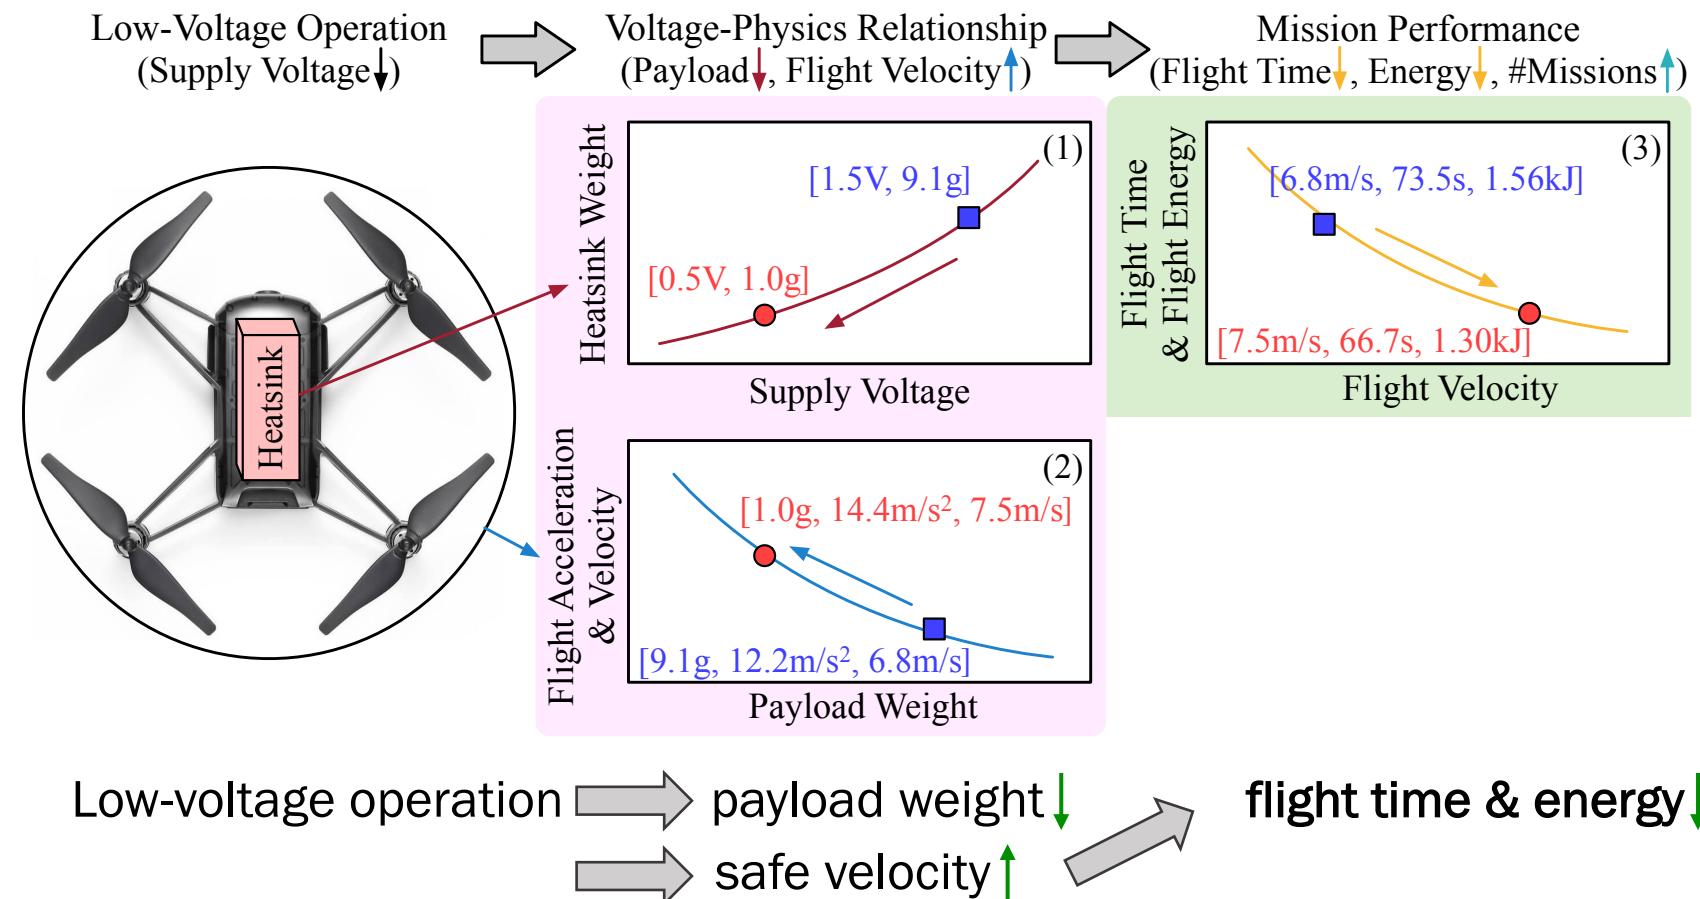
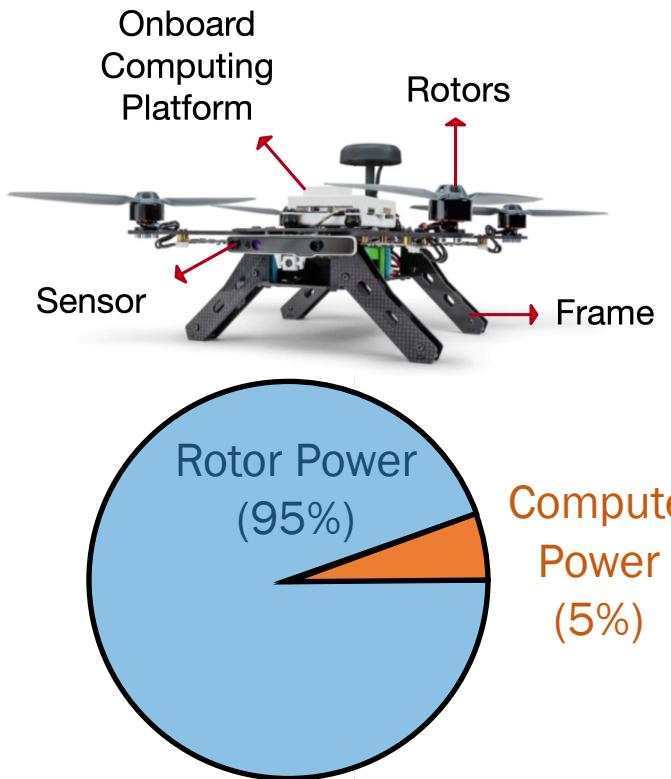
[2] Celsia Heatsink Size Simulator



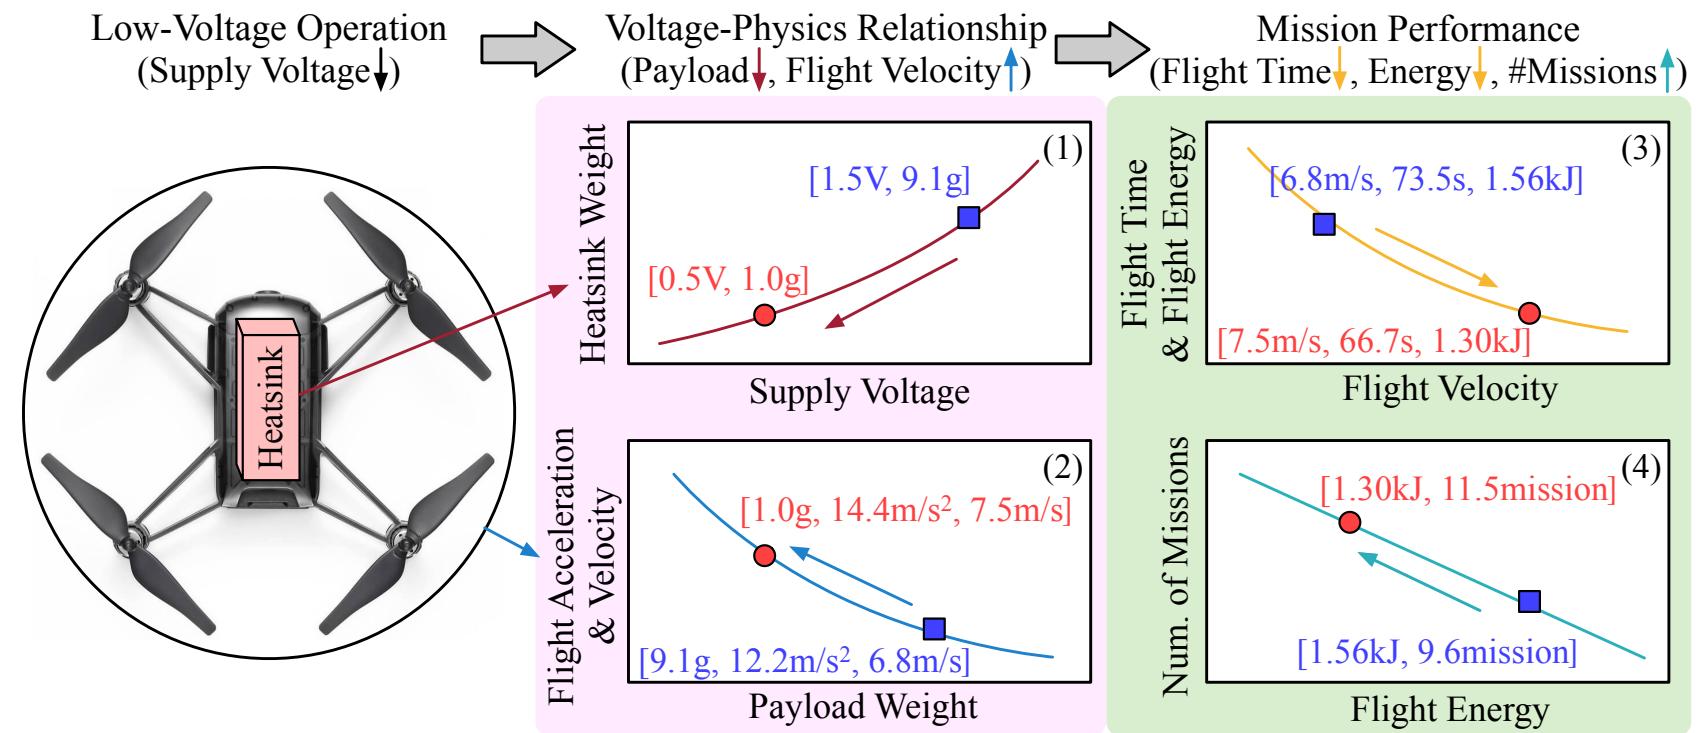
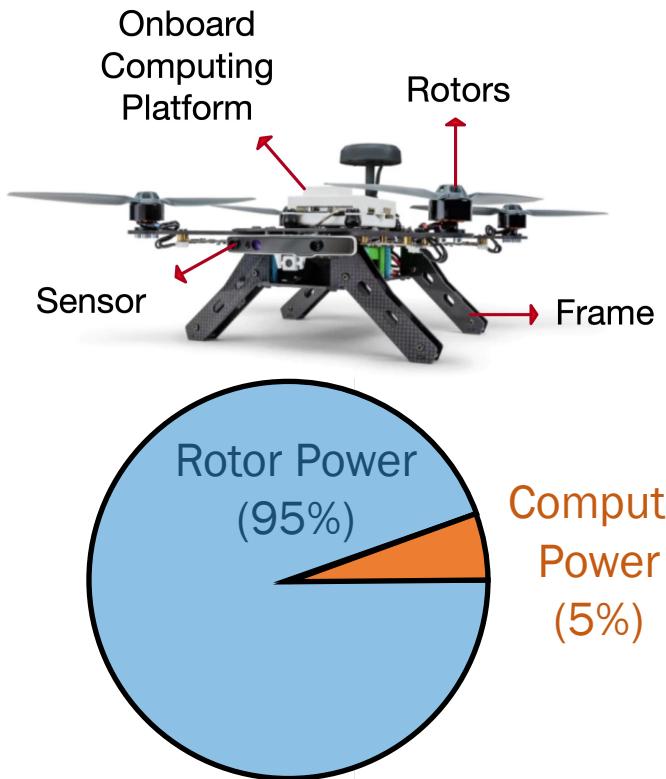
Small Compute Power, Large System Impact!



Small Compute Power, Large System Impact!



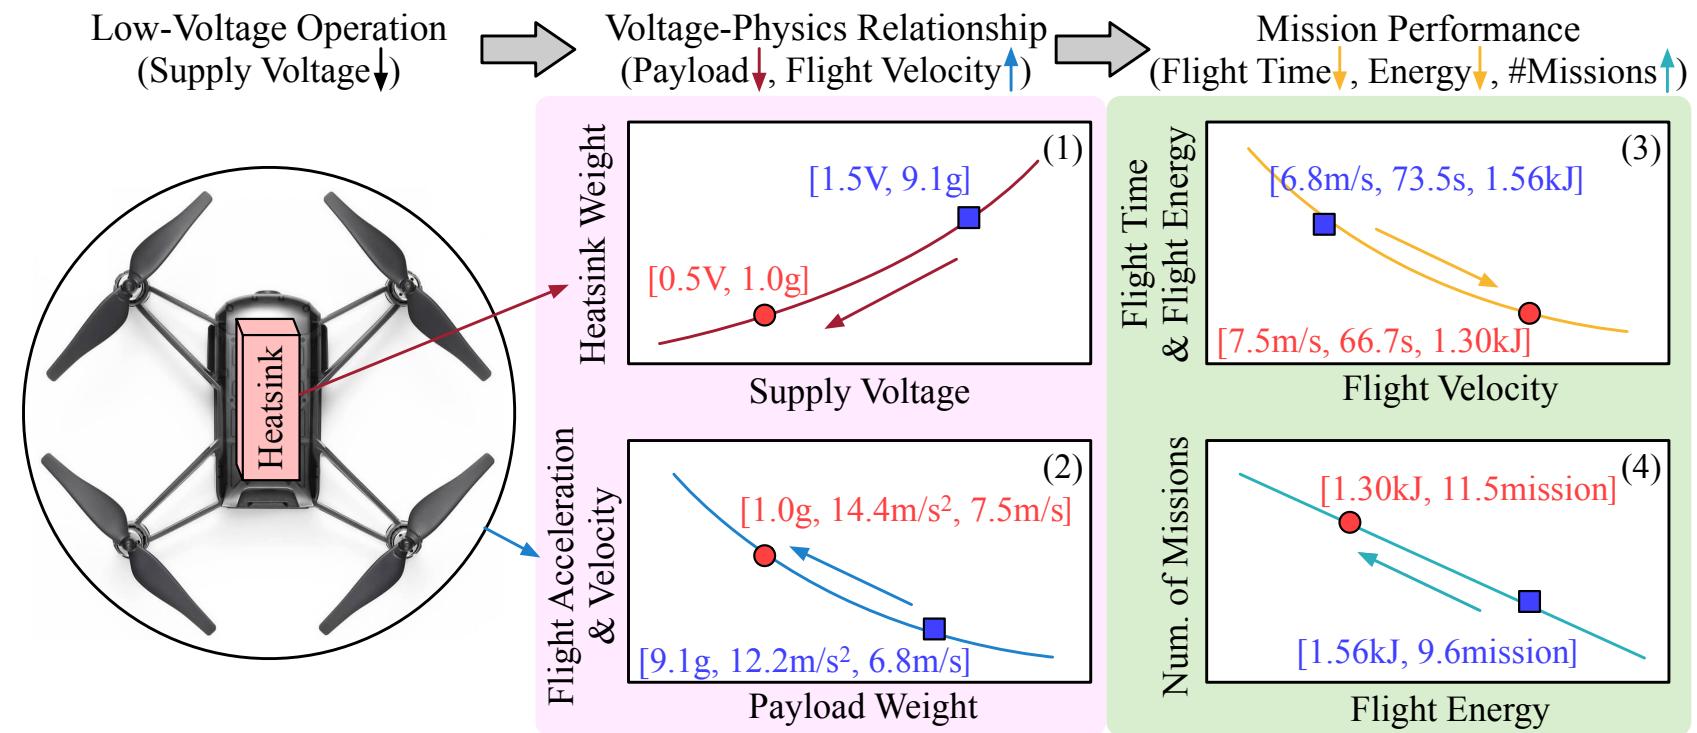
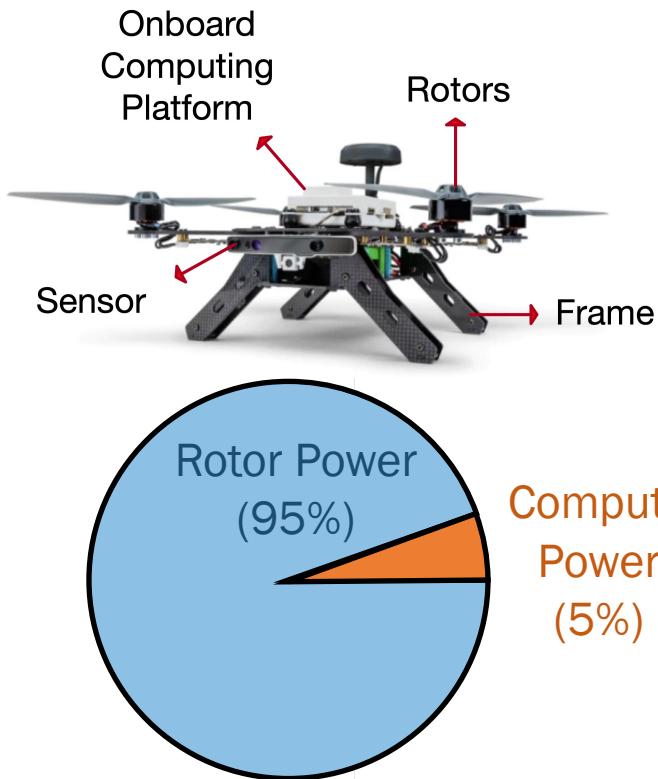
Small Compute Power, Large System Impact!



Low-voltage operation \rightarrow payload weight \downarrow
 \rightarrow safe velocity \uparrow \rightarrow flight time & energy \downarrow
num of missions \uparrow



Small Compute Power, Large System Impact!

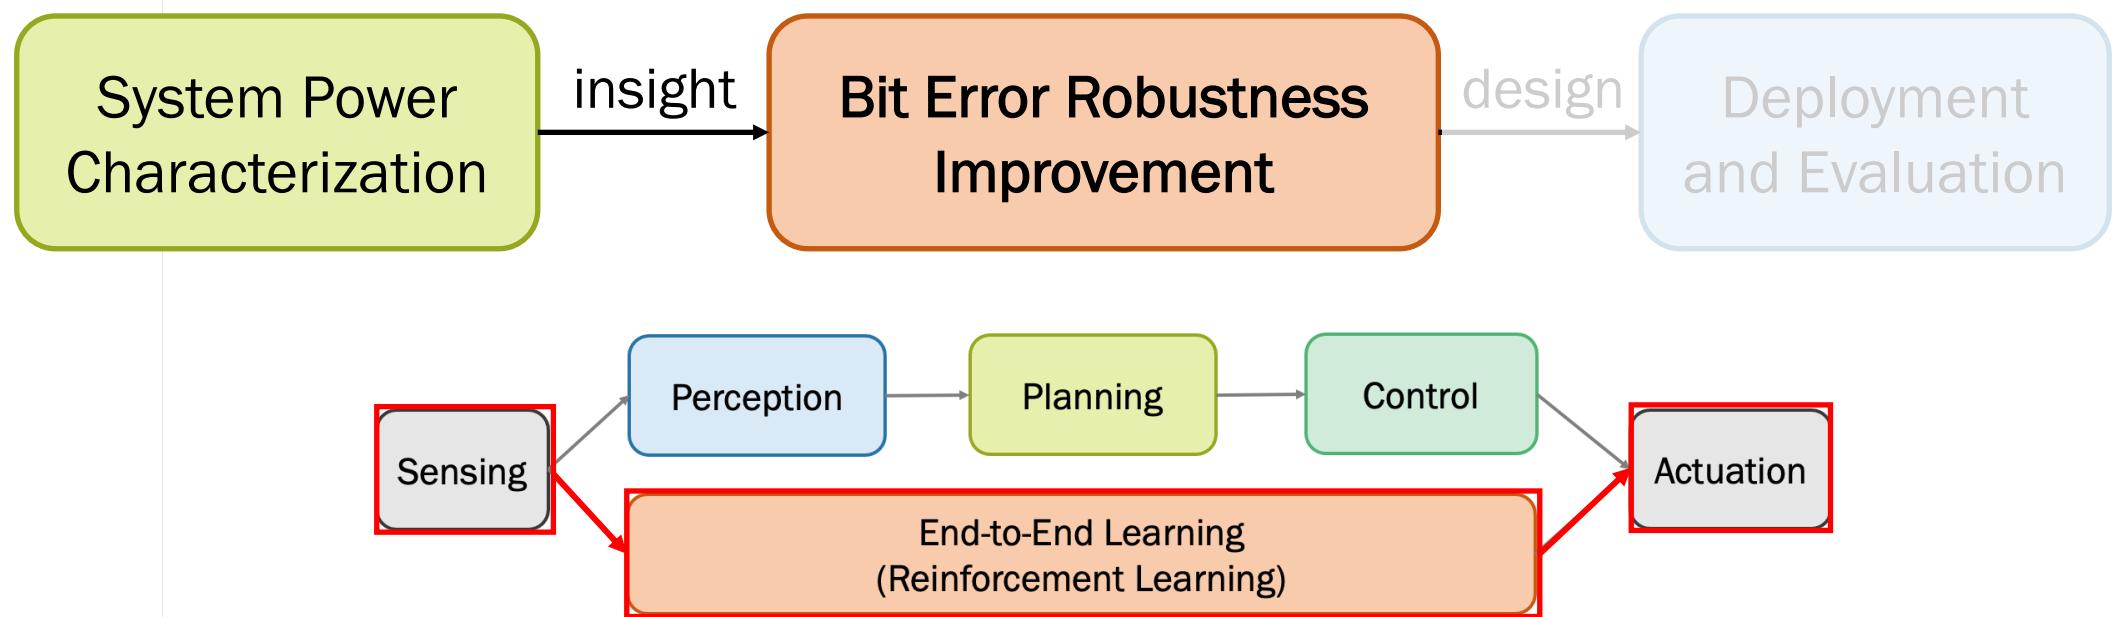


Compute power has huge impacts on end-to-end autonomous system mission energy

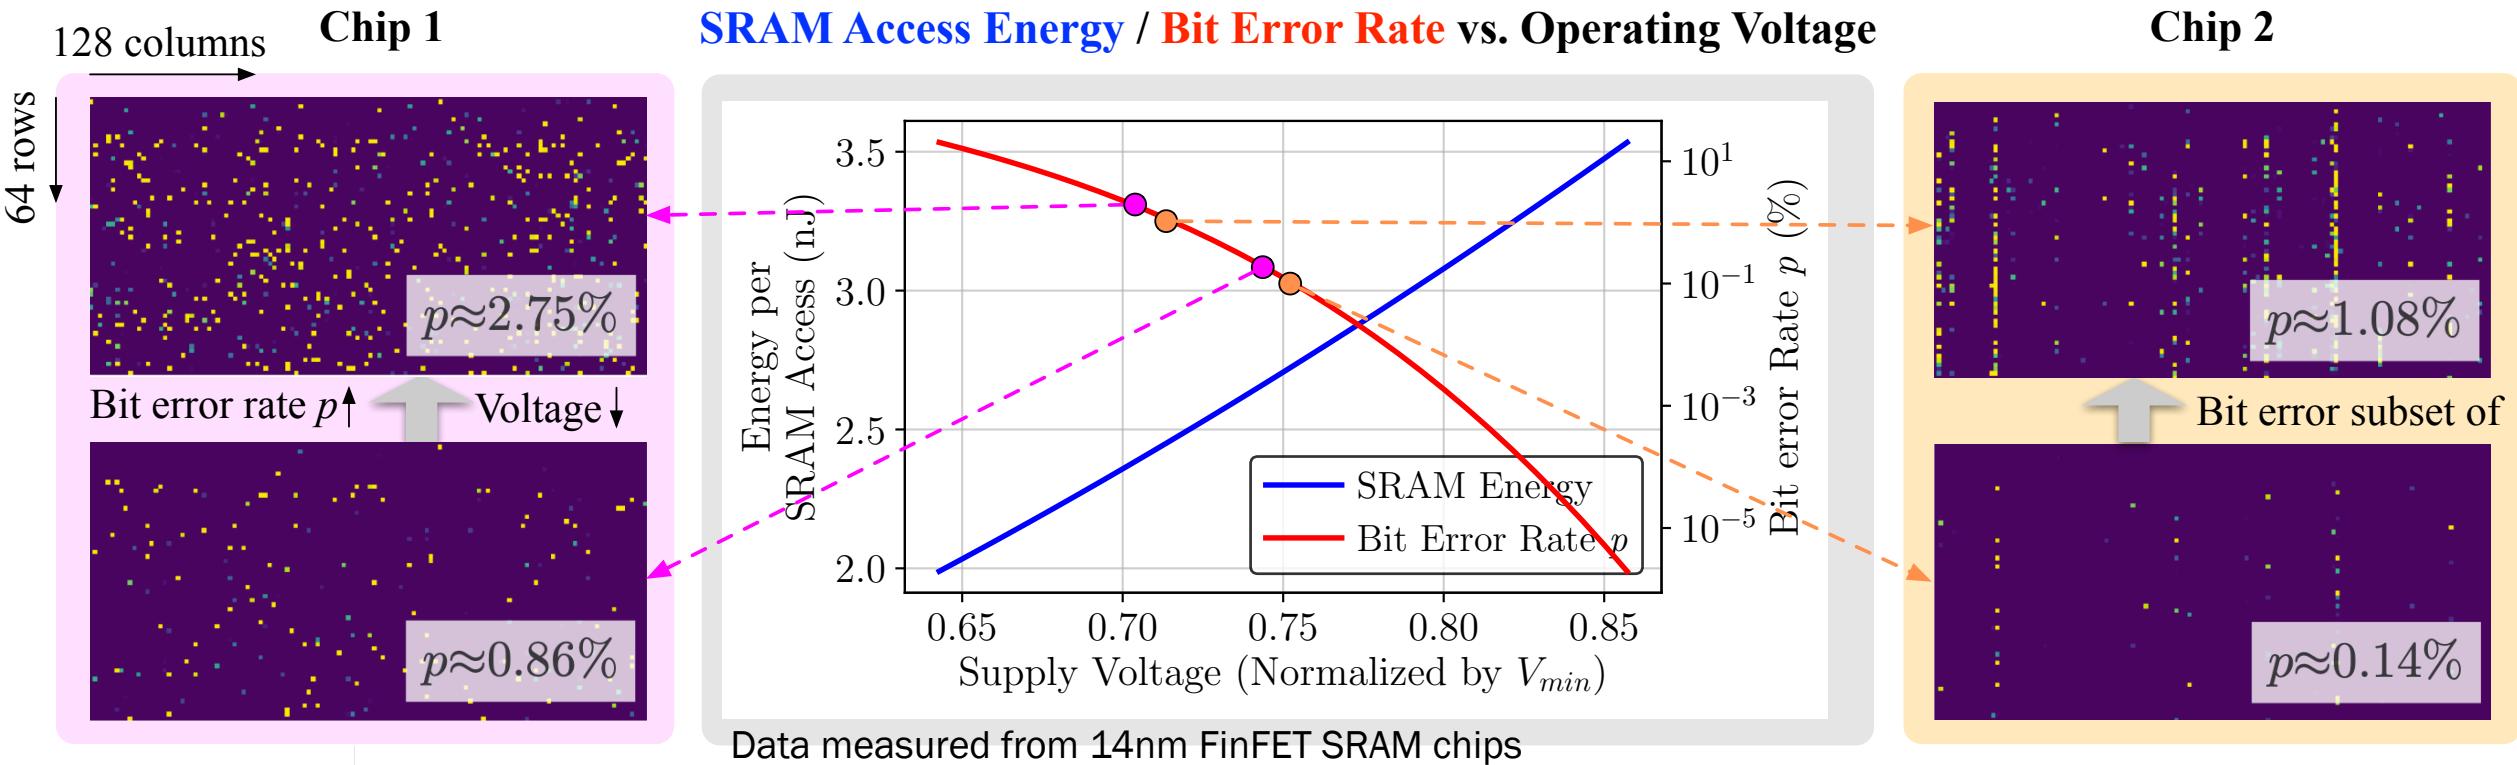


BERRY Framework

(BERRY: Bit Error Robustness for Energy-Efficient Reinforcement Learning-Based Autonomous Systems)



Challenge: Low Voltage Induces Faults



Operating below rated voltage range results in memory bit errors, negatively impacting safety



BERRY Framework

- Design Principle: Cross-layer robust learning framework, integrates *algorithm-level error-aware training* with *hardware-level* bit failures, aiming to optimize system-level **robustness/efficiency/performance** under **low operating voltage**.



BERRY Framework

- Design Principle: Cross-layer robust learning framework, integrates *algorithm-level error-aware training* with *hardware-level* bit failures, aiming to optimize system-level *robustness/efficiency/performance* under *low operating voltage*.

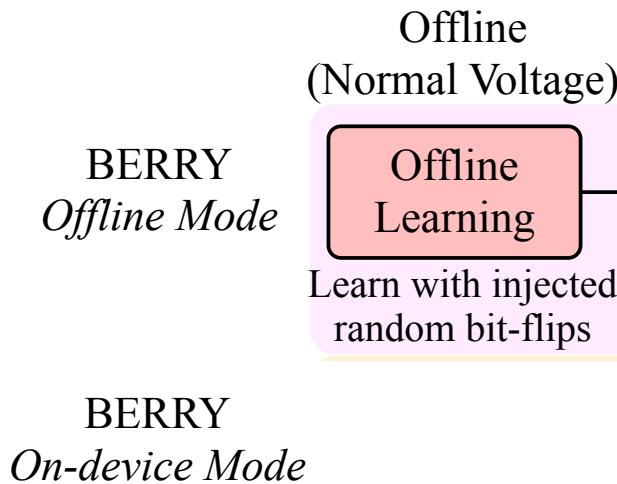
BERRY
Offline Mode

BERRY
On-device Mode



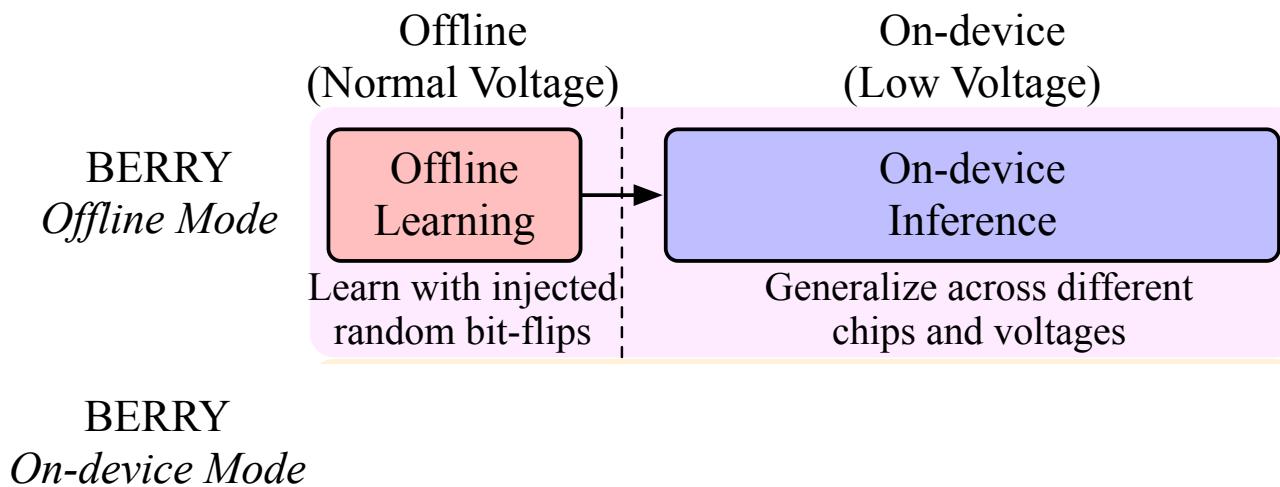
BERRY Framework

- Design Principle: Cross-layer robust learning framework, integrates *algorithm-level error-aware training* with *hardware-level* bit failures, aiming to optimize system-level **robustness/efficiency/performance** under **low operating voltage**.



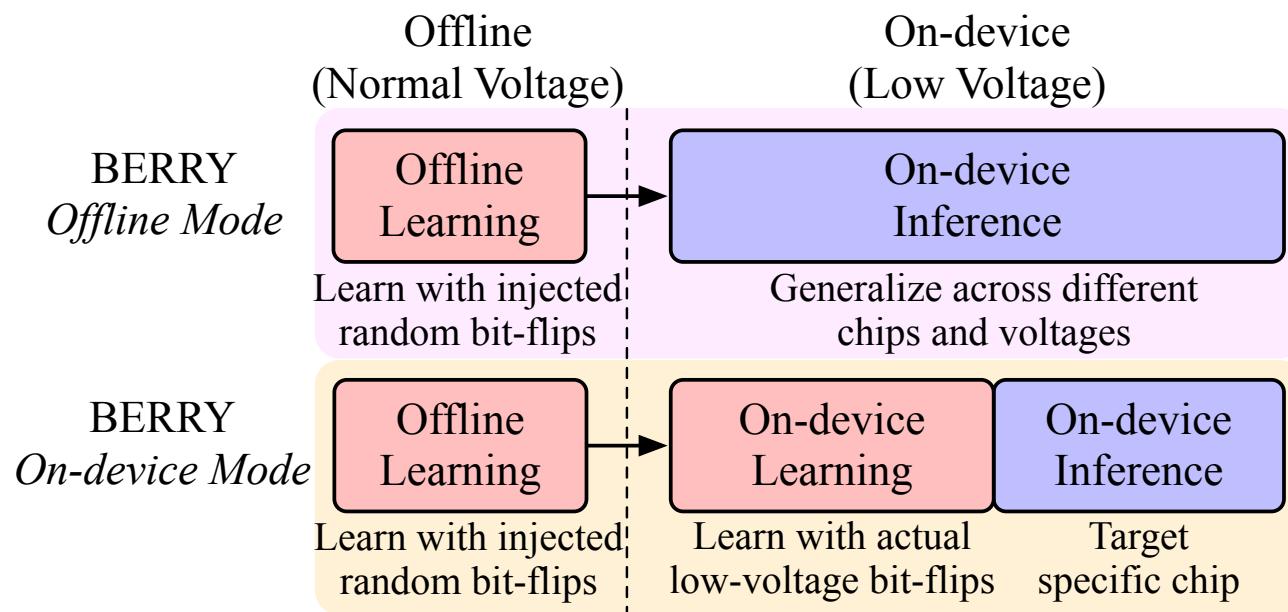
BERRY Framework

- **Design Principle:** Cross-layer robust learning framework, integrates *algorithm-level error-aware training* with *hardware-level* bit failures, aiming to optimize system-level **robustness/efficiency/performance** under **low operating voltage**.



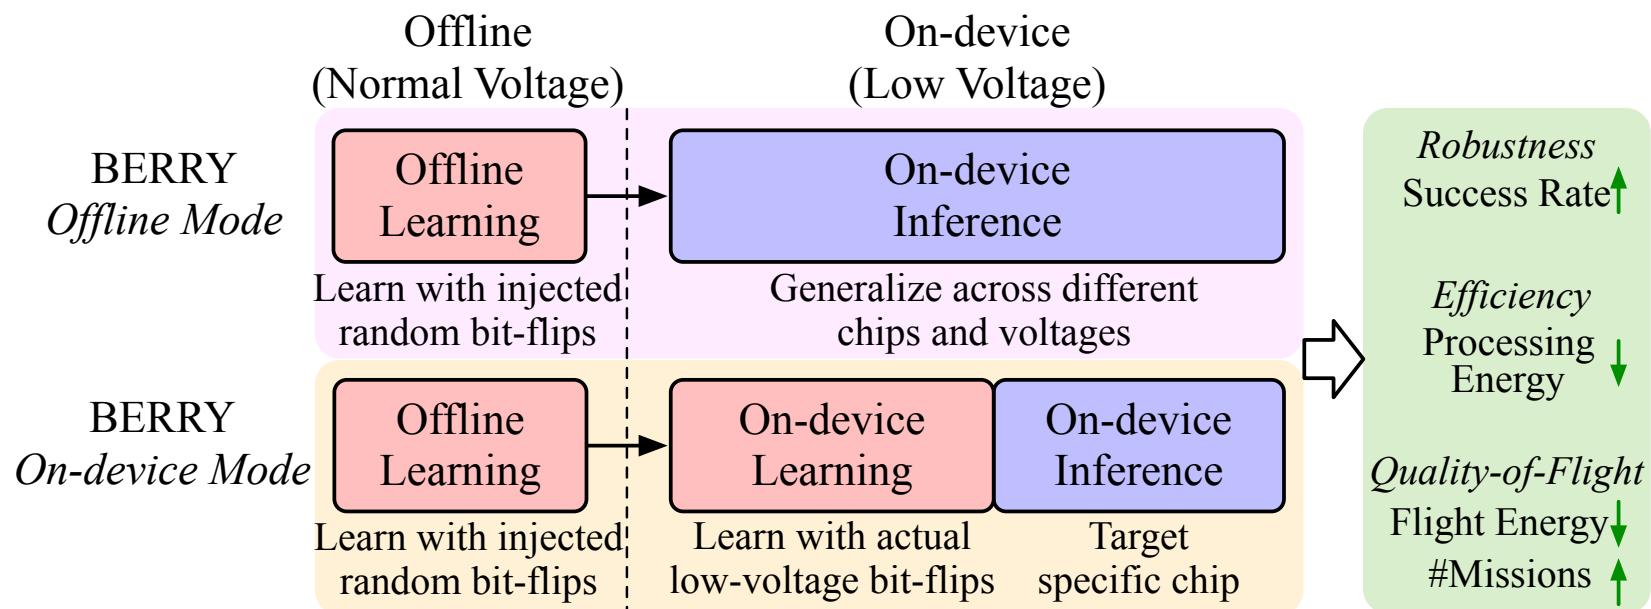
BERRY Framework

- **Design Principle:** Cross-layer robust learning framework, integrates *algorithm-level error-aware training* with *hardware-level* bit failures, aiming to optimize *system-level robustness/efficiency/performance* under *low operating voltage*.



BERRY Framework

- **Design Principle:** Cross-layer robust learning framework, integrates *algorithm-level error-aware training* with *hardware-level* bit failures, aiming to optimize system-level **robustness/efficiency/performance** under **low operating voltage**.



BERRY Framework

Algorithm 1 BERRY Robust Error-Aware Training Framework for Reinforcement Learning-Based Autonomous Systems

```
1: procedure BERRY( $p$ )
2:   Initialize action-value function  $Q$  with weight  $\theta$ 
3:   Initialize target action-value function  $\hat{Q}$  with weight  $\theta^- = \theta$ 
4:   for episode  $e = 1$  to  $E$  do
5:     for time step  $t = 1$  to  $T$  do
6:       Given state  $s_t$ , take action  $a_t$  based on  $Q$  ( $\epsilon$ -greedy)
7:       Obtain reward  $r_t$  and reach new state  $s_{t+1}$ 
8:       Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$ 
9:       // Experience replay
10:      Sample a mini-batch  $\{(s_j, a_j, r_j, s_{j+1})\}_{b=1}^B$  from  $D$ 
11:      // Clean training pass
12:      Set  $y_j = r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)})$ 
13:       $\Delta^{(t)} = \nabla_\theta \sum_{b=1}^B (Q(s_j, a_j; \theta^{(t)}) - y_j)^2$ 
14:      // Perturbed training pass, inject bit errors at rate p
15:       $\tilde{\theta}^{(t)} = BErr_p(\theta^{(t)})$ 
16:      Set  $\tilde{y}_j = (r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \tilde{\theta}^{-(t)}))$ 
17:       $\tilde{\Delta}^{(t)} = \nabla_\theta \sum_{b=1}^B (Q(s_j, a_j; \tilde{\theta}^{(t)}) - \tilde{y}_j)^2$ 
18:      // Average gradients and update w.r.t  $\theta$ 
19:       $\theta^{(t+1)} = \theta^{(t)} - \alpha(\Delta^{(t)} + \tilde{\Delta}^{(t)})$ 
20:      // Periodic update of target network
21:      Every  $C$  steps reset  $\hat{Q} = Q$ , i.e., set  $\theta^- = \theta$ 
22:    end for
23:  end for
24: Output: Bit-error robust action-value function  $Q(\theta)$ 
```



BERRY Framework

Start: Initialize policy

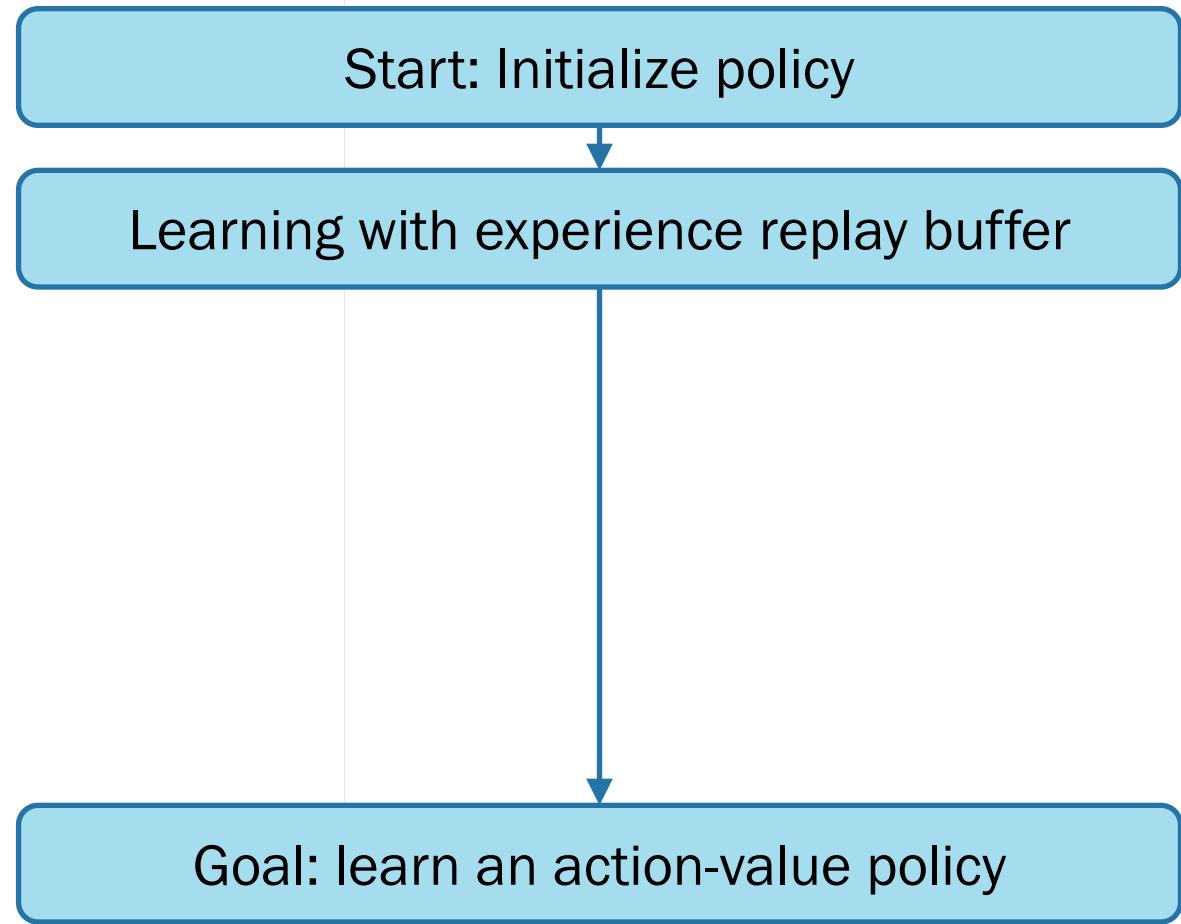
Goal: learn an action-value policy

Algorithm 1 BERRY Robust Error-Aware Training Framework for Reinforcement Learning-Based Autonomous Systems

```
1: procedure BERRY( $p$ )
2:   Initialize action-value function  $Q$  with weight  $\theta$ 
3:   Initialize target action-value function  $\hat{Q}$  with weight  $\theta^- = \theta$ 
4:   for episode  $e = 1$  to  $E$  do
5:     for time step  $t = 1$  to  $T$  do
6:       Given state  $s_t$ , take action  $a_t$  based on  $Q$  ( $\epsilon$ -greedy)
7:       Obtain reward  $r_t$  and reach new state  $s_{t+1}$ 
8:       Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$ 
9:       // Experience replay
10:      Sample a mini-batch  $\{(s_j, a_j, r_j, s_{j+1})\}_{b=1}^B$  from  $D$ 
11:      // Clean training pass
12:      Set  $y_j = r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)})$ 
13:       $\Delta^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \theta^{(t)}) - y_j)^2$ 
14:      // Perturbed training pass, inject bit errors at rate p
15:       $\tilde{\theta}^{(t)} = BErr_p(\theta^{(t)})$ 
16:      Set  $\tilde{y}_j = (r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \tilde{\theta}^{-(t)}))$ 
17:       $\tilde{\Delta}^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \tilde{\theta}^{(t)}) - \tilde{y}_j)^2$ 
18:      // Average gradients and update w.r.t  $\theta$ 
19:       $\theta^{(t+1)} = \theta^{(t)} - \alpha(\Delta^{(t)} + \tilde{\Delta}^{(t)})$ 
20:      // Periodic update of target network
21:      Every  $C$  steps reset  $\hat{Q} = Q$ , i.e., set  $\theta^- = \theta$ 
22:    end for
23:  end for
24: Output: Bit-error robust action-value function  $Q(\theta)$ 
```



BERRY Framework



Algorithm 1 BERRY Robust Error-Aware Training Framework for Reinforcement Learning-Based Autonomous Systems

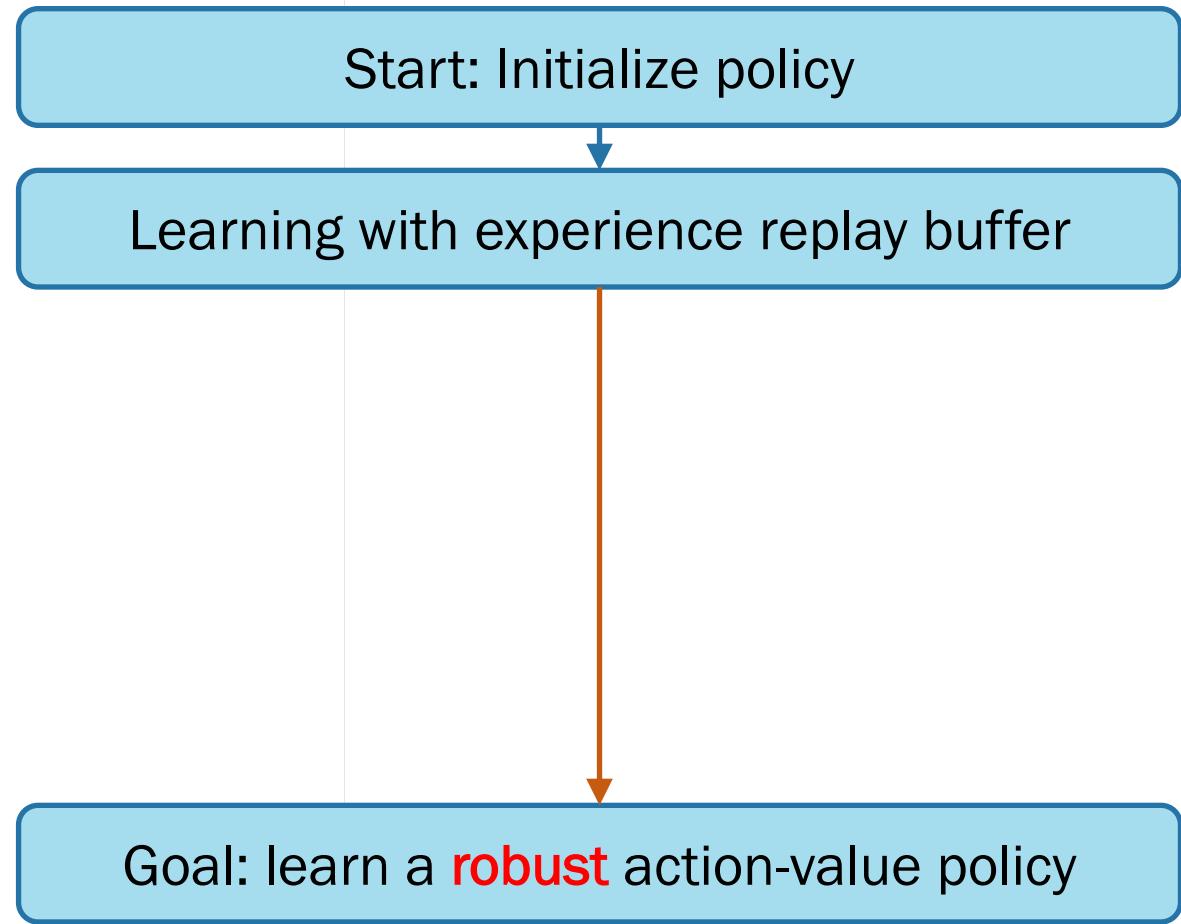
```

1: procedure BERRY( $p$ )
2:   Initialize action-value function  $Q$  with weight  $\theta$ 
3:   Initialize target action-value function  $\hat{Q}$  with weight  $\theta^- = \theta$ 
4:   for episode  $e = 1$  to  $E$  do
5:     for time step  $t = 1$  to  $T$  do
6:       Given state  $s_t$ , take action  $a_t$  based on  $Q$  ( $\epsilon$ -greedy)
7:       Obtain reward  $r_t$  and reach new state  $s_{t+1}$ 
8:       Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$ 
9:       // Experience replay
10:      Sample a mini-batch  $\{(s_j, a_j, r_j, s_{j+1})\}_{b=1}^B$  from  $D$ 
11:      // Clean training pass
12:      Set  $y_j = r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)})$ 
13:       $\Delta^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \theta^{(t)}) - y_j)^2$ 
14:      // Perturbed training pass, inject bit errors at rate  $p$ 
15:       $\tilde{\theta}^{(t)} = BErr_p(\theta^{(t)})$ 
16:      Set  $\tilde{y}_j = (r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \tilde{\theta}^{-(t)}))$ 
17:       $\tilde{\Delta}^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \tilde{\theta}^{(t)}) - \tilde{y}_j)^2$ 
18:      // Average gradients and update w.r.t  $\theta$ 
19:       $\theta^{(t+1)} = \theta^{(t)} - \alpha(\Delta^{(t)} + \tilde{\Delta}^{(t)})$ 
20:      // Periodic update of target network
21:      Every  $C$  steps reset  $\hat{Q} = Q$ , i.e., set  $\theta^- = \theta$ 
22:    end for
23:  end for
24: Output: Bit-error robust action-value function  $Q(\theta)$ 

```



BERRY Framework

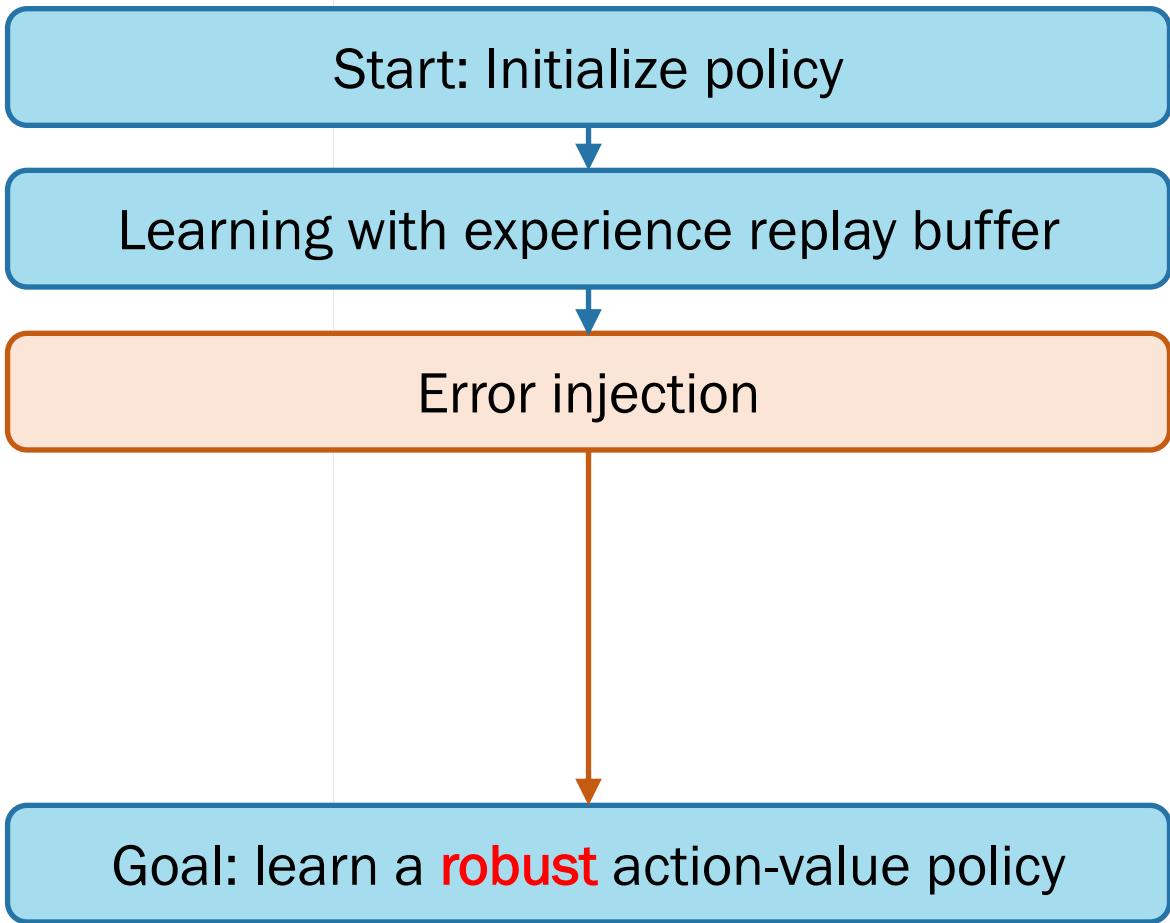


Algorithm 1 BERRY Robust Error-Aware Training Framework for Reinforcement Learning-Based Autonomous Systems

```
1: procedure BERRY( $p$ )
2:   Initialize action-value function  $Q$  with weight  $\theta$ 
3:   Initialize target action-value function  $\hat{Q}$  with weight  $\theta^- = \theta$ 
4:   for episode  $e = 1$  to  $E$  do
5:     for time step  $t = 1$  to  $T$  do
6:       Given state  $s_t$ , take action  $a_t$  based on  $Q$  ( $\epsilon$ -greedy)
7:       Obtain reward  $r_t$  and reach new state  $s_{t+1}$ 
8:       Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$ 
9:       // Experience replay
10:      Sample a mini-batch  $\{(s_j, a_j, r_j, s_{j+1})\}_{b=1}^B$  from  $D$ 
11:      // Clean training pass
12:      Set  $y_j = r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)})$ 
13:       $\Delta^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \theta^{(t)}) - y_j)^2$ 
14:      // Perturbed training pass, inject bit errors at rate  $p$ 
15:       $\tilde{\theta}^{(t)} = BErr_p(\theta^{(t)})$ 
16:      Set  $\tilde{y}_j = (r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \tilde{\theta}^{-(t)}))$ 
17:       $\tilde{\Delta}^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \tilde{\theta}^{(t)}) - \tilde{y}_j)^2$ 
18:      // Average gradients and update w.r.t  $\theta$ 
19:       $\theta^{(t+1)} = \theta^{(t)} - \alpha(\Delta^{(t)} + \tilde{\Delta}^{(t)})$ 
20:      // Periodic update of target network
21:      Every  $C$  steps reset  $\hat{Q} = Q$ , i.e., set  $\theta^- = \theta$ 
22:    end for
23:  end for
24: Output: Bit-error robust action-value function  $Q(\theta)$ 
```



BERRY Framework



Algorithm 1 BERRY Robust Error-Aware Training Framework for Reinforcement Learning-Based Autonomous Systems

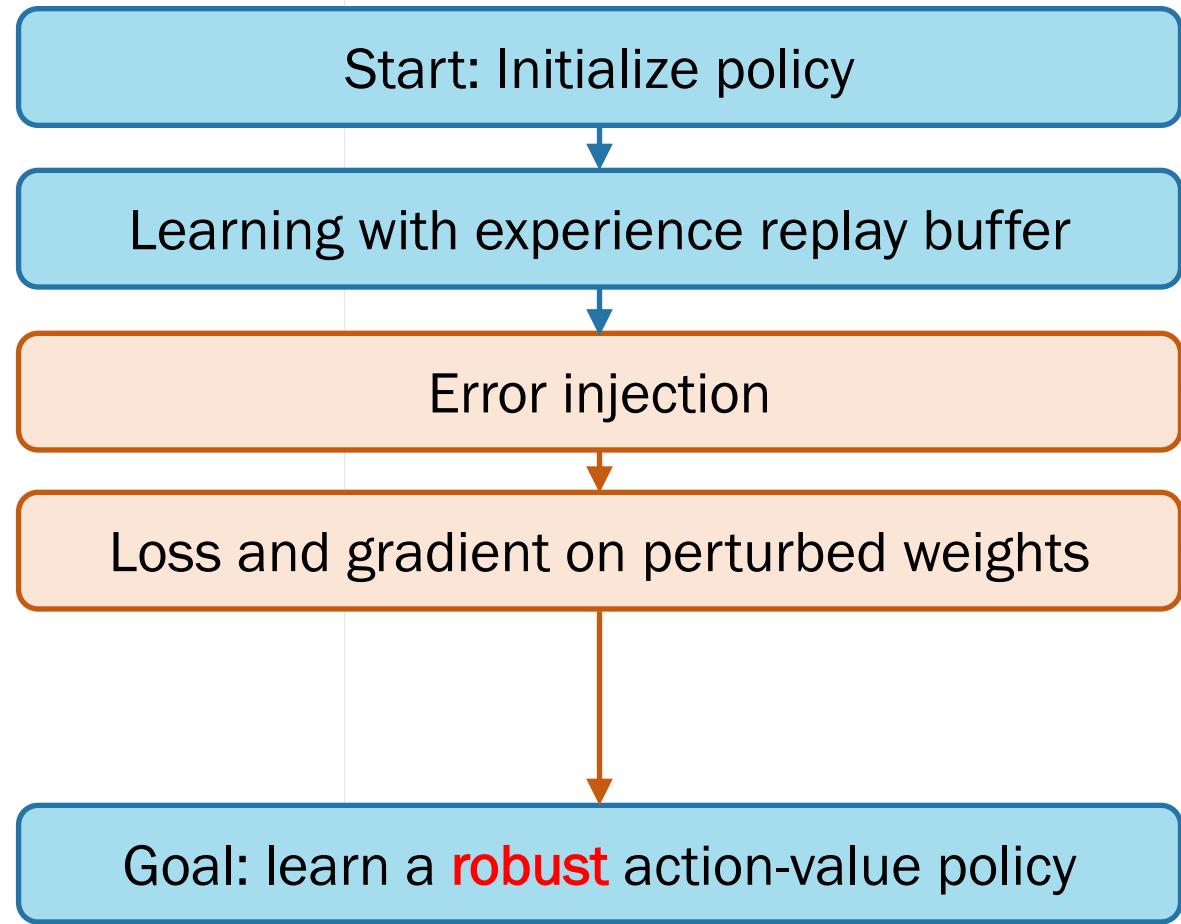
```

1: procedure BERRY( $p$ )
2:   Initialize action-value function  $Q$  with weight  $\theta$ 
3:   Initialize target action-value function  $\hat{Q}$  with weight  $\theta^- = \theta$ 
4:   for episode  $e = 1$  to  $E$  do
5:     for time step  $t = 1$  to  $T$  do
6:       Given state  $s_t$ , take action  $a_t$  based on  $Q$  ( $\epsilon$ -greedy)
7:       Obtain reward  $r_t$  and reach new state  $s_{t+1}$ 
8:       Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$ 
9:       // Experience replay
10:      Sample a mini-batch  $\{(s_j, a_j, r_j, s_{j+1})\}_{b=1}^B$  from  $D$ 
11:      // Clean training pass
12:      Set  $y_j = r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)})$ 
13:       $\Delta^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \theta^{(t)}) - y_j)^2$ 
14:      // Perturbed training pass, inject bit errors at rate p
15:       $\tilde{\theta}^{(t)} = BErr_p(\theta^{(t)})$ 
16:      Set  $\tilde{y}_j = (r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)}))$ 
17:       $\tilde{\Delta}^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \tilde{\theta}^{(t)}) - \tilde{y}_j)^2$ 
18:      // Average gradients and update w.r.t  $\theta$ 
19:       $\theta^{(t+1)} = \theta^{(t)} - \alpha(\Delta^{(t)} + \tilde{\Delta}^{(t)})$ 
20:      // Periodic update of target network
21:      Every  $C$  steps reset  $\hat{Q} = Q$ , i.e., set  $\theta^- = \theta$ 
22:    end for
23:  end for
24: Output: Bit-error robust action-value function  $Q(\theta)$ 

```



BERRY Framework



Algorithm 1 BERRY Robust Error-Aware Training Framework for Reinforcement Learning-Based Autonomous Systems

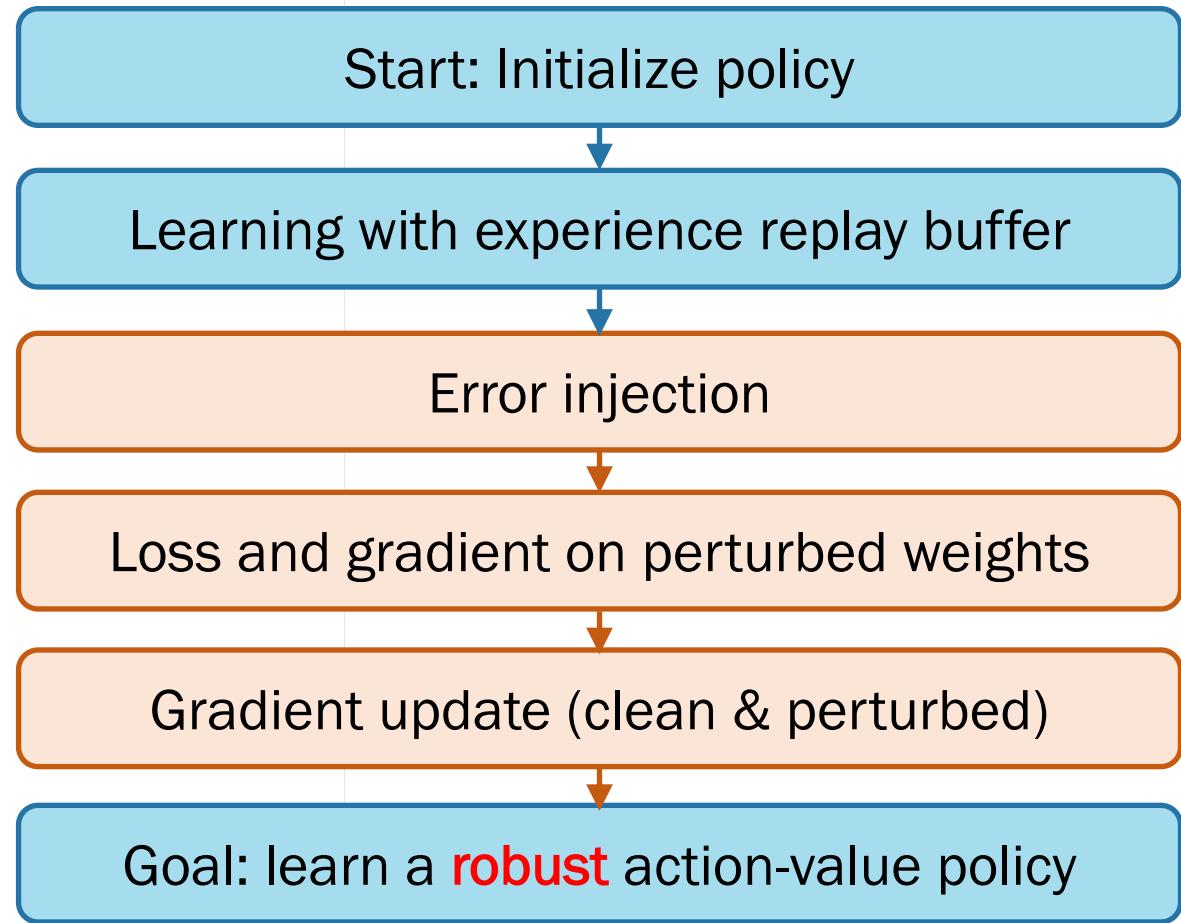
```

1: procedure BERRY( $p$ )
2:   Initialize action-value function  $Q$  with weight  $\theta$ 
3:   Initialize target action-value function  $\hat{Q}$  with weight  $\theta^- = \theta$ 
4:   for episode  $e = 1$  to  $E$  do
5:     for time step  $t = 1$  to  $T$  do
6:       Given state  $s_t$ , take action  $a_t$  based on  $Q$  ( $\epsilon$ -greedy)
7:       Obtain reward  $r_t$  and reach new state  $s_{t+1}$ 
8:       Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$ 
9:       // Experience replay
10:      Sample a mini-batch  $\{(s_j, a_j, r_j, s_{j+1})\}_{b=1}^B$  from  $D$ 
11:      // Clean training pass
12:      Set  $y_j = r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)})$ 
13:       $\Delta^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \theta^{(t)}) - y_j)^2$ 
14:      // Perturbed training pass, inject bit errors at rate  $p$ 
15:       $\tilde{\theta}^{(t)} = BErr_p(\theta^{(t)})$ 
16:      Set  $\tilde{y}_j = (r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)}))$ 
17:       $\tilde{\Delta}^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \tilde{\theta}^{(t)}) - \tilde{y}_j)^2$ 
18:      // Average gradients and update w.r.t  $\theta$ 
19:       $\theta^{(t+1)} = \theta^{(t)} - \alpha(\Delta^{(t)} + \tilde{\Delta}^{(t)})$ 
20:      // Periodic update of target network
21:      Every  $C$  steps reset  $\hat{Q} = Q$ , i.e., set  $\theta^- = \theta$ 
22:    end for
23:  end for
24: Output: Bit-error robust action-value function  $Q(\theta)$ 

```



BERRY Framework



Algorithm 1 BERRY Robust Error-Aware Training Framework for Reinforcement Learning-Based Autonomous Systems

```

1: procedure BERRY( $p$ )
2:   Initialize action-value function  $Q$  with weight  $\theta$ 
3:   Initialize target action-value function  $\hat{Q}$  with weight  $\theta^- = \theta$ 
4:   for episode  $e = 1$  to  $E$  do
5:     for time step  $t = 1$  to  $T$  do
6:       Given state  $s_t$ , take action  $a_t$  based on  $Q$  ( $\epsilon$ -greedy)
7:       Obtain reward  $r_t$  and reach new state  $s_{t+1}$ 
8:       Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$ 
9:       // Experience replay
10:      Sample a mini-batch  $\{(s_j, a_j, r_j, s_{j+1})\}_{b=1}^B$  from  $D$ 
11:      // Clean training pass
12:      Set  $y_j = r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)})$ 
13:       $\Delta^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \theta^{(t)}) - y_j)^2$ 
14:      // Perturbed training pass, inject bit errors at rate  $p$ 
15:       $\tilde{\theta}^{(t)} = BErr_p(\theta^{(t)})$ 
16:      Set  $\tilde{y}_j = (r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta^{-(t)}))$ 
17:       $\tilde{\Delta}^{(t)} = \nabla_{\theta} \sum_{b=1}^B (Q(s_j, a_j; \tilde{\theta}^{(t)}) - \tilde{y}_j)^2$ 
18:      // Average gradients and update w.r.t  $\theta$ 
19:       $\theta^{(t+1)} = \theta^{(t)} - \alpha(\Delta^{(t)} + \tilde{\Delta}^{(t)})$ 
20:      // Periodic update of target network
21:      Every  $C$  steps reset  $\hat{Q} = Q$ , i.e., set  $\theta^- = \theta$ 
22:    end for
23:  end for
24: Output: Bit-error robust action-value function  $Q(\theta)$ 

```

*Please refer to paper for algorithm details



BERRY Framework Highlights

Previous Error-Aware Training

- Profiled errors injected during offline training, resulting in robust model during inference
- Supervised learning tasks (e.g., object classification)
- Model robustness

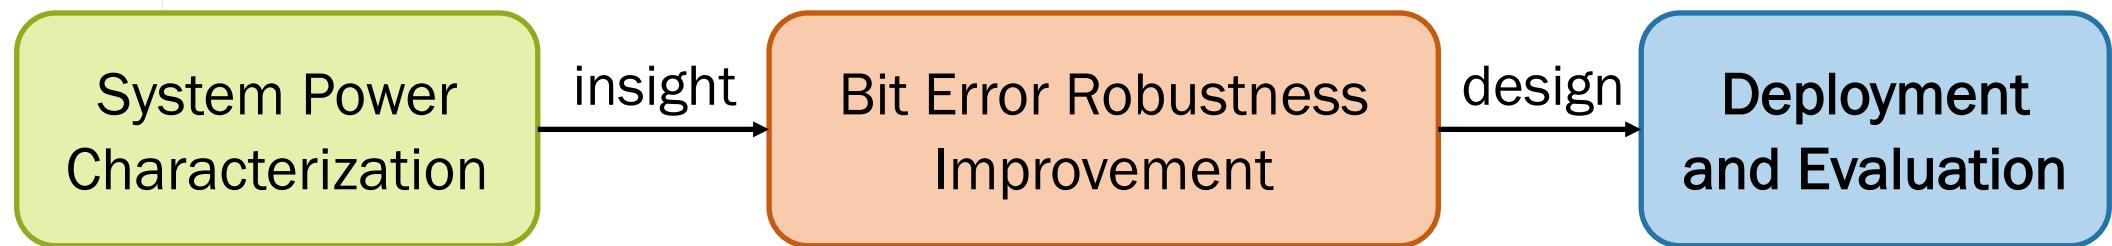
BERRY

- BERRY targets **learning (both offline and on-device)**, meaning that learning can occur on low-voltage devices with bit errors affecting parameters
- BERRY targets **reinforcement learning**
- BERRY tackles complex relationship between low-voltage compute and **cyber-physical** autonomous systems



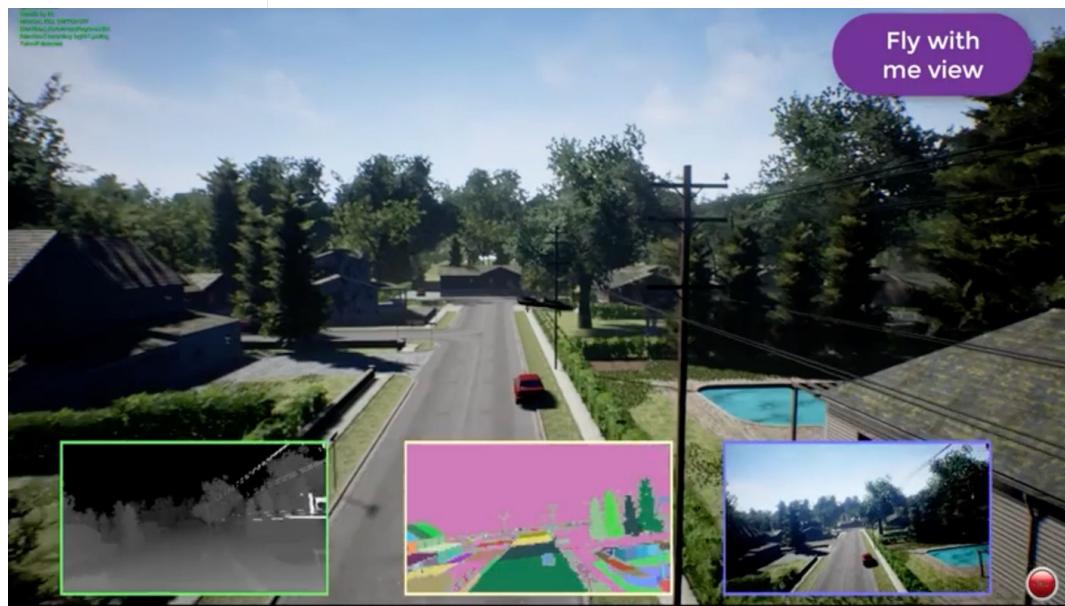
BERRY Framework

(BERRY: Bit Error Robustness for Energy-Efficient Reinforcement Learning-Based Autonomous Systems)



Experimental Setup

- Simulation Platform

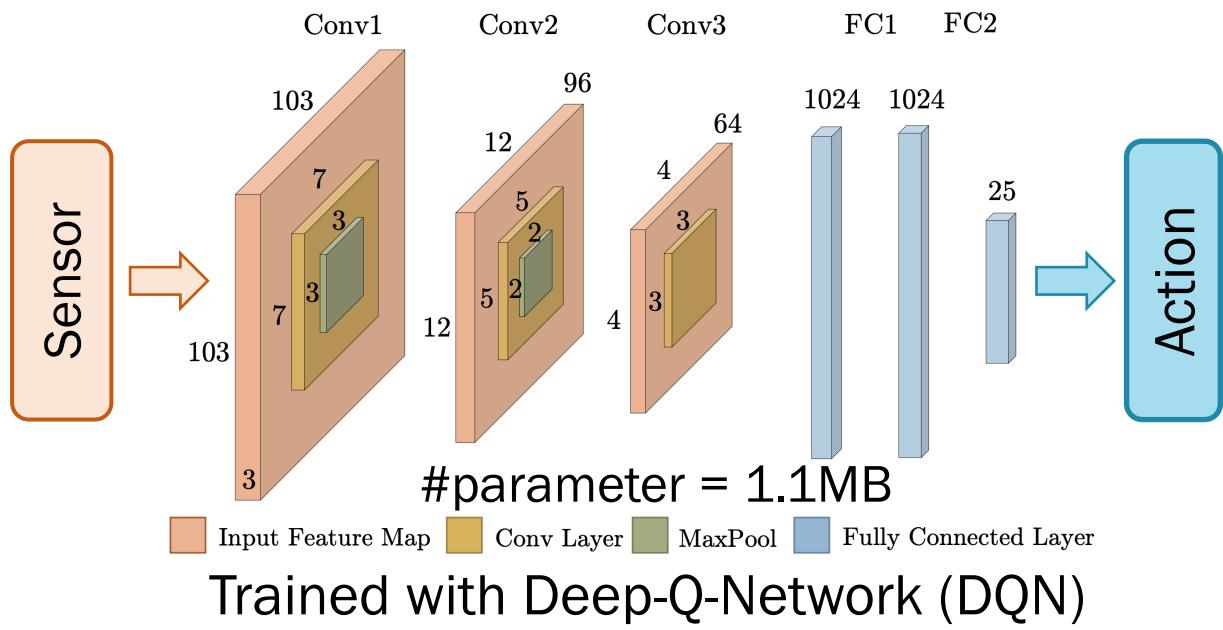


Unreal Engine + AirSim
(3D realistic environments) (Drone dynamics)

- Task

- point-to-point autonomous navigation (e.g., package delivery)

- Policy Architecture



Experimental Setup

- Drone Platforms
 - Bitcraze Crazyflie



- DJI Tello



Nano-Drone
27g takeoff weight
15g max payload
250mAh battery

Micro-Drone
80g takeoff weight
70g max payload
1100mAh battery

- Evaluation Metrics

- Compute-level: Processing energy

- System-level:

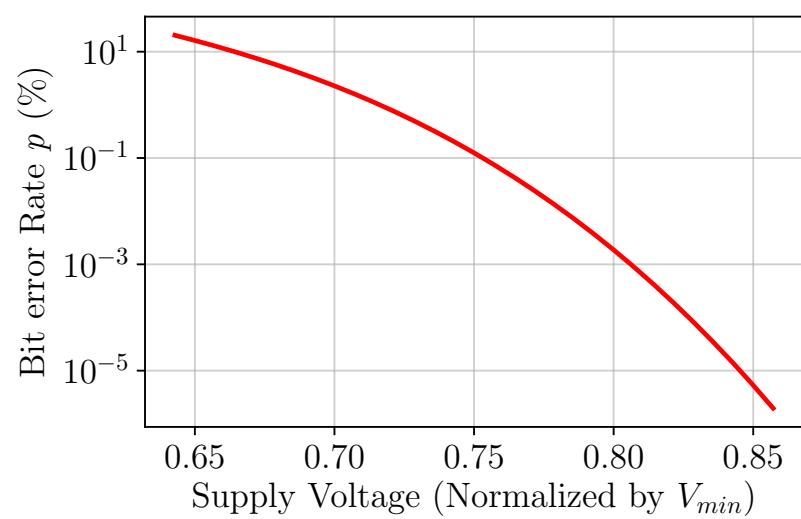
- Flight success rate ↑
- Flight time ↓
- Flight energy ↓
- Number of completed missions ↑

All reported results are averaged from 500 different fault maps



Robustness Improvement

Bit Error Rate vs. Voltage

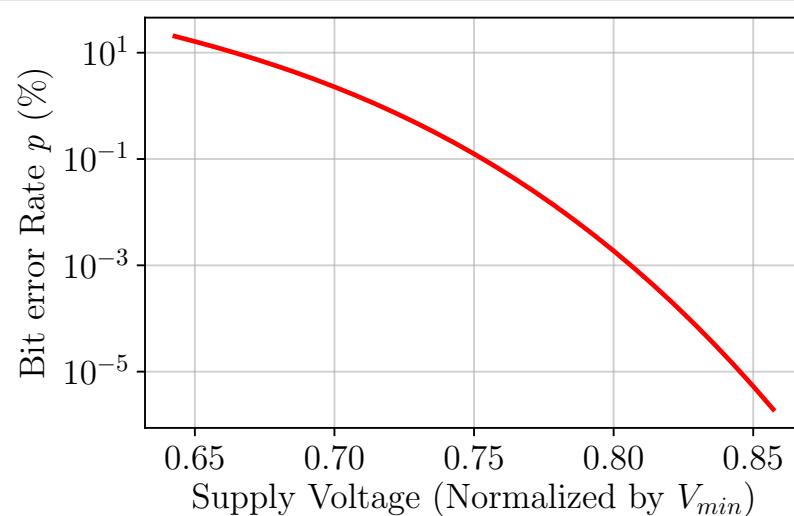


Voltage ↓ → Bit error rate ↑



Robustness Improvement

Bit Error Rate vs. Voltage



Voltage ↓ → Bit error rate ↑

Flight Success Rate under Various Bit Error Rates p (Low Voltages)

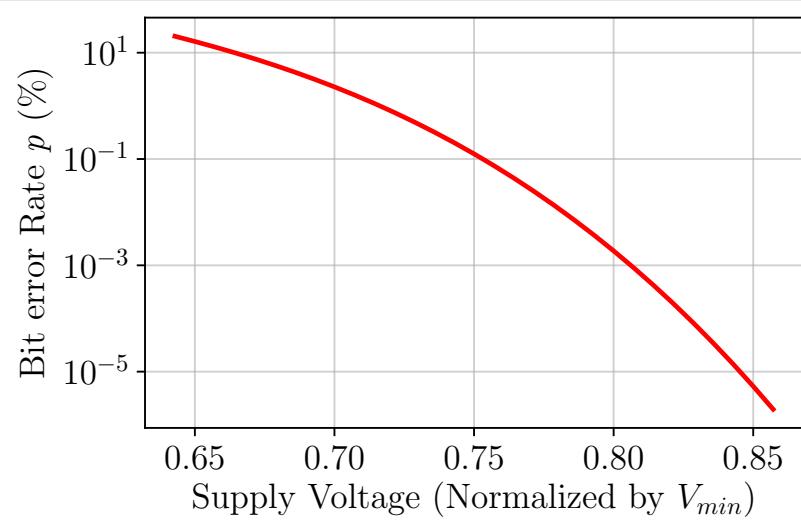
Autonomy Scheme	Success Rate	Bit-Error Success Rate			
		$p=0.01\%$	$p=0.05\%$	$p=0.1\%$	$p=0.5\%$
Classical	88.4%	84.0%	78.2%	69.2%	48.6%

(Success Rate: percentage of successful flight of all flight trails)



Robustness Improvement

Bit Error Rate vs. Voltage



Voltage ↓ → Bit error rate ↑

Flight Success Rate under Various Bit Error Rates p (Low Voltages)

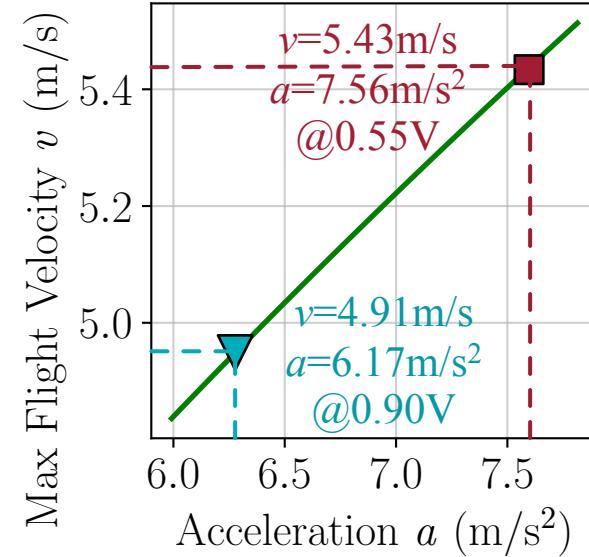
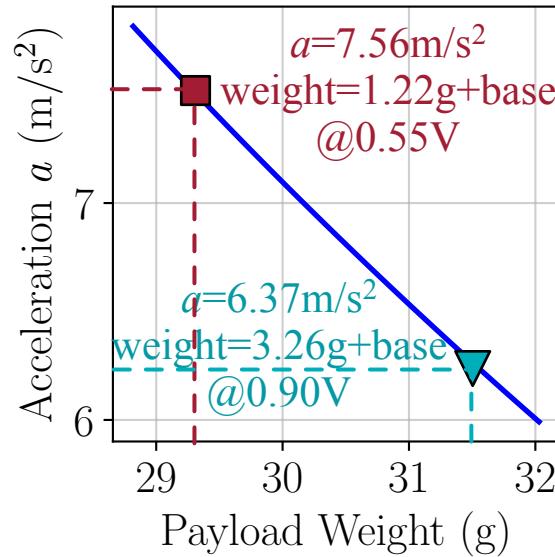
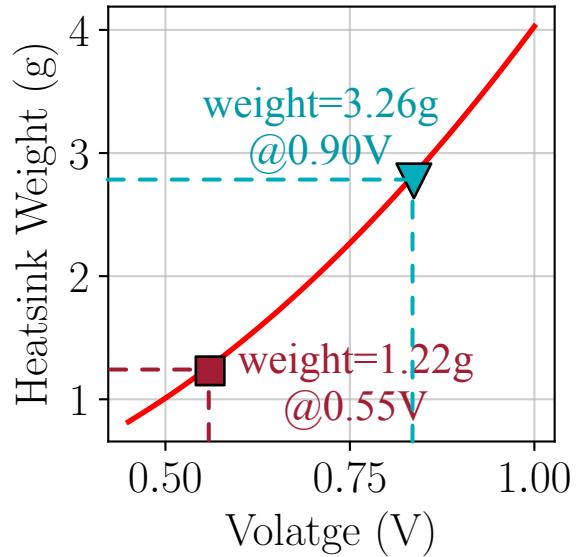
Autonomy Scheme	Success Rate	Bit-Error Success Rate			
		$p=0.01\%$	$p=0.05\%$	$p=0.1\%$	$p=0.5\%$
Classical	88.4%	84.0%	78.2%	69.2%	48.6%
BERRY	88.8%	88.6%	86.6%	84.4%	79.8%

(Success Rate: percentage of successful flight of all flight trails)

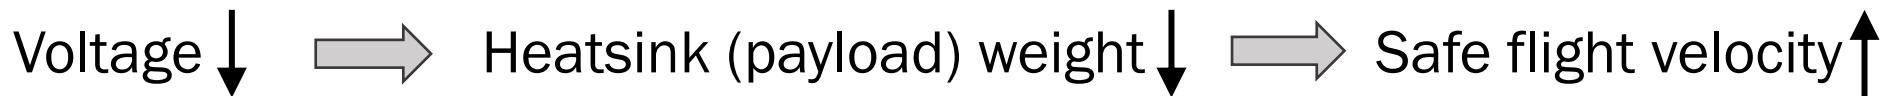
BERRY improves mission success rate under bit failures



Mission Efficiency Improvement



(Characterized from Bitcraze Crazyfile drone configuration)



Lowering operating voltage can bring benefits to cyber-physical drones (e.g., safe velocity)



Mission Efficiency Improvement

Low-Voltage Operation	
Voltage (V)	Bit Error Rate p (%)
1	0
$0.86V_{min}$	1.96×10^{-6}
$0.84V_{min}$	1.38×10^{-5}
$0.83V_{min}$	8.23×10^{-5}
$0.81V_{min}$	4.22×10^{-4}
$0.80V_{min}$	1.87×10^{-3}
$0.79V_{min}$	7.25×10^{-3}
$0.77V_{min}$	2.47×10^{-2}
$0.76V_{min}$	7.49×10^{-2}
$0.74V_{min}$	2.03×10^{-1}
$0.73V_{min}$	4.98×10^{-1}
$0.71V_{min}$	1.11
$0.68V_{min}$	5.80
$0.64V_{min}$	20.36

Bit errors increase



Mission Efficiency Improvement

Low-Voltage Operation	
Voltage (V)	Bit Error Rate p (%)
1	0
$0.86V_{min}$	1.96×10^{-6}
$0.84V_{min}$	1.38×10^{-5}
$0.83V_{min}$	8.23×10^{-5}
$0.81V_{min}$	4.22×10^{-4}
$0.80V_{min}$	1.87×10^{-3}
$0.79V_{min}$	7.25×10^{-3}
$0.77V_{min}$	2.47×10^{-2}
$0.76V_{min}$	7.49×10^{-2}
$0.74V_{min}$	2.03×10^{-1}
$0.73V_{min}$	4.98×10^{-1}
$0.71V_{min}$	1.11
$0.68V_{min}$	5.80
$0.64V_{min}$	20.36

Bit errors increase



Mission Efficiency Improvement

Low-Voltage Operation		Processing
Voltage (V)	Bit Error Rate p (%)	Energy Savings (%)
1	0	-
$0.86V_{min}$	1.96×10^{-6}	2.77×
$0.84V_{min}$	1.38×10^{-5}	2.87×
$0.83V_{min}$	8.23×10^{-5}	2.97×
$0.81V_{min}$	4.22×10^{-4}	3.07×
$0.80V_{min}$	1.87×10^{-3}	3.18×
$0.79V_{min}$	7.25×10^{-3}	3.30×
$0.77V_{min}$	2.47×10^{-2}	3.43×
$0.76V_{min}$	7.49×10^{-2}	3.55×
$0.74V_{min}$	2.03×10^{-1}	3.69×
$0.73V_{min}$	4.98×10^{-1}	3.84×
$0.71V_{min}$	1.11	3.99×
$0.68V_{min}$	5.80	4.42×
$0.64V_{min}$	20.36	4.93×

Bit errors increase
Save processing energy



Mission Efficiency Improvement

Low-Voltage Operation		Processing	Robustness	
Voltage (V)	Bit Error Rate p (%)	Energy Savings (%)	Success Rate (%)	
1	0	-	88.4	
$0.86V_{min}$	1.96×10^{-6}	2.77 \times	88.0	
$0.84V_{min}$	1.38×10^{-5}	2.87 \times	89.2	
$0.83V_{min}$	8.23×10^{-5}	2.97 \times	89.0	
$0.81V_{min}$	4.22×10^{-4}	3.07 \times	88.8	
$0.80V_{min}$	1.87×10^{-3}	3.18 \times	88.6	
$0.79V_{min}$	7.25×10^{-3}	3.30 \times	88.6	
$0.77V_{min}$	2.47×10^{-2}	3.43\times	88.4	
$0.76V_{min}$	7.49×10^{-2}	3.55 \times	86.2	
$0.74V_{min}$	2.03×10^{-1}	3.69 \times	83.4	
$0.73V_{min}$	4.98×10^{-1}	3.84 \times	79.0	
$0.71V_{min}$	1.11	3.99 \times	74.4	
$0.68V_{min}$	5.80	4.42 \times	63.2	
$0.64V_{min}$	20.36	4.93 \times	50.4	

Bit errors increase

Save processing energy

Maintain robustness under faults



Mission Efficiency Improvement

Low-Voltage Operation		Processing	Robustness	Autonomous System Mission-Level Quality-of-Flight	
Voltage (V)	Bit Error Rate p (%)	Energy Savings (%)	Success Rate (%)	Flight Distance (m)	
1	0	-	88.4	14.89	
$0.86V_{min}$	1.96×10^{-6}	2.77 \times	88.0	14.93	
$0.84V_{min}$	1.38×10^{-5}	2.87 \times	89.2	14.86	
$0.83V_{min}$	8.23×10^{-5}	2.97 \times	89.0	14.91	
$0.81V_{min}$	4.22×10^{-4}	3.07 \times	88.8	14.96	
$0.80V_{min}$	1.87×10^{-3}	3.18 \times	88.6	14.94	
$0.79V_{min}$	7.25×10^{-3}	3.30 \times	88.6	14.94	
$0.77V_{min}$	2.47×10^{-2}	3.43\times	88.4	14.91	
$0.76V_{min}$	7.49×10^{-2}	3.55 \times	86.2	15.71	
$0.74V_{min}$	2.03×10^{-1}	3.69 \times	83.4	16.58	
$0.73V_{min}$	4.98×10^{-1}	3.84 \times	79.0	18.03	
$0.71V_{min}$	1.11	3.99 \times	74.4	19.46	
$0.68V_{min}$	5.80	4.42 \times	63.2	21.84	
$0.64V_{min}$	20.36	4.93 \times	50.4	24.52	

Bit errors increase

Save processing energy

Maintain robustness under faults

Maintain trajectory distance under faults



Mission Efficiency Improvement

Low-Voltage Operation		Processing	Robustness	Autonomous System Mission-Level Quality-of-Flight		
Voltage (V)	Bit Error Rate p (%)	Energy Savings (%)	Success Rate (%)	Flight Distance (m)	Flight Time (s)	
1	0	-	88.4	14.89	6.81	
$0.86V_{min}$	1.96×10^{-6}	2.77 \times	88.0	14.93	6.51	
$0.84V_{min}$	1.38×10^{-5}	2.87 \times	89.2	14.86	6.48	
$0.83V_{min}$	8.23×10^{-5}	2.97 \times	89.0	14.91	6.46	
$0.81V_{min}$	4.22×10^{-4}	3.07 \times	88.8	14.96	6.45	
$0.80V_{min}$	1.87×10^{-3}	3.18 \times	88.6	14.94	6.42	
$0.79V_{min}$	7.25×10^{-3}	3.30 \times	88.6	14.94	6.39	
$0.77V_{min}$	2.47×10^{-2}	3.43\times	88.4	14.91	6.35	
$0.76V_{min}$	7.49×10^{-2}	3.55 \times	86.2	15.71	6.67	
$0.74V_{min}$	2.03×10^{-1}	3.69 \times	83.4	16.58	7.03	
$0.73V_{min}$	4.98×10^{-1}	3.84 \times	79.0	18.03	7.61	
$0.71V_{min}$	1.11	3.99 \times	74.4	19.46	8.18	
$0.68V_{min}$	5.80	4.42 \times	63.2	21.84	9.09	
$0.64V_{min}$	20.36	4.93 \times	50.4	24.52	10.11	

BERRY improves mission efficiency (flight time↓,



Mission Efficiency Improvement

Low-Voltage Operation		Processing	Robustness	Autonomous System Mission-Level Quality-of-Flight			
Voltage (V)	Bit Error Rate p (%)	Energy Savings (%)	Success Rate (%)	Flight Distance (m)	Flight Time (s)	Flight Energy E_{flight} (J)	E_{flight} Savings
1	0	-	88.4	14.89	6.81	53.19	-
$0.86V_{min}$	1.96×10^{-6}	2.77 \times	88.0	14.93	6.51	47.23	-11.21%
$0.84V_{min}$	1.38×10^{-5}	2.87 \times	89.2	14.86	6.48	46.66	-12.28%
$0.83V_{min}$	8.23×10^{-5}	2.97 \times	89.0	14.91	6.46	46.41	-12.73%
$0.81V_{min}$	4.22×10^{-4}	3.07 \times	88.8	14.96	6.45	46.22	-13.11%
$0.80V_{min}$	1.87×10^{-3}	3.18 \times	88.6	14.94	6.42	45.80	-13.90%
$0.79V_{min}$	7.25×10^{-3}	3.30 \times	88.6	14.94	6.39	45.38	-14.67%
$0.77V_{min}$	2.47×10^{-2}	3.43\times	88.4	14.91	6.35	44.88	-15.62%
$0.76V_{min}$	7.49×10^{-2}	3.55 \times	86.2	15.71	6.67	46.90	-11.82%
$0.74V_{min}$	2.03×10^{-1}	3.69 \times	83.4	16.58	7.03	49.14	-7.61%
$0.73V_{min}$	4.98×10^{-1}	3.84 \times	79.0	18.03	7.61	52.98	-0.39%
$0.71V_{min}$	1.11	3.99 \times	74.4	19.46	8.18	56.62	-6.45%
$0.68V_{min}$	5.80	4.42 \times	63.2	21.84	9.09	61.96	+16.49%
$0.64V_{min}$	20.36	4.93 \times	50.4	24.52	10.11	67.83	+27.53%

BERRY improves mission efficiency (flight time ↓, flight energy ↓, (15.62%)



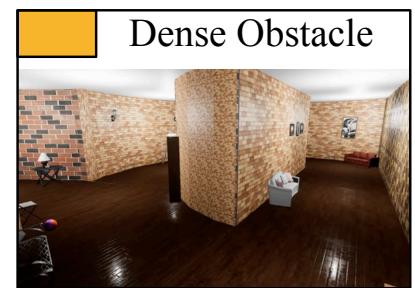
Mission Efficiency Improvement

Low-Voltage Operation		Processing	Robustness	Autonomous System			Mission-Level Quality-of-Flight		
Voltage (V)	Bit Error Rate p (%)	Energy Savings (%)	Success Rate (%)	Flight Distance (m)	Flight Time (s)	Flight Energy E_{flight} (J)	E_{flight} Savings	Num. of Missions $N_{mission}$	$N_{mission}$ Improvements
1	0	-	88.4	14.89	6.81	53.19	-	55.35	-
$0.86V_{min}$	1.96×10^{-6}	2.77 \times	88.0	14.93	6.51	47.23	-11.21%	62.05	+12.12%
$0.84V_{min}$	1.38×10^{-5}	2.87 \times	89.2	14.86	6.48	46.66	-12.28%	63.66	+15.03%
$0.83V_{min}$	8.23×10^{-5}	2.97 \times	89.0	14.91	6.46	46.41	-12.73%	63.85	+15.37%
$0.81V_{min}$	4.22×10^{-4}	3.07 \times	88.8	14.96	6.45	46.22	-13.11%	63.98	+15.61%
$0.80V_{min}$	1.87×10^{-3}	3.18 \times	88.6	14.94	6.42	45.80	-13.90%	64.42	+16.40%
$0.79V_{min}$	7.25×10^{-3}	3.30 \times	88.6	14.94	6.39	45.38	-14.67%	65.01	+17.46%
$0.77V_{min}$	2.47×10^{-2}	3.43\times	88.4	14.91	6.35	44.88	-15.62%	65.59	+18.51%
$0.76V_{min}$	7.49×10^{-2}	3.55 \times	86.2	15.71	6.67	46.90	-11.82%	61.20	+10.58%
$0.74V_{min}$	2.03×10^{-1}	3.69 \times	83.4	16.58	7.03	49.14	-7.61%	56.52	+2.12%
$0.73V_{min}$	4.98×10^{-1}	3.84 \times	79.0	18.03	7.61	52.98	-0.39%	49.66	-10.27%
$0.71V_{min}$	1.11	3.99 \times	74.4	19.46	8.18	56.62	-6.45%	43.75	-20.95%
$0.68V_{min}$	5.80	4.42 \times	63.2	21.84	9.09	61.96	+16.49%	33.96	-38.64%
$0.64V_{min}$	20.36	4.93 \times	50.4	24.52	10.11	67.83	+27.53%	24.74	-55.30%

BERRY improves mission efficiency (flight time ↓, flight energy ↓, number of completed missions ↑)
 (15.62%) (18.51%)



Effectiveness Across Environments



Environments	Sparse Obstacle	Medium Obstacle	Dense Obstacle
Operating Voltage			
Flight Energy Savings			
Num. of Success Mission Increase			



Effectiveness Across Environments



Environments	Sparse Obstacle	Medium Obstacle	Dense Obstacle
Operating Voltage	$0.76 V_{min}$	$0.77 V_{min}$	$0.80 V_{min}$
Flight Energy Savings			
Num. of Success Mission Increase			



Effectiveness Across Environments



Environments	Sparse Obstacle	Medium Obstacle	Dense Obstacle
Operating Voltage	$0.76 V_{min}$	$0.77 V_{min}$	$0.80 V_{min}$
Flight Energy Savings	14.0%	15.6%	15.6%
Num. of Success Mission Increase			

Effectiveness Across Environments



Environments	Sparse Obstacle	Medium Obstacle	Dense Obstacle
Operating Voltage	$0.76 V_{min}$	$0.77 V_{min}$	$0.80 V_{min}$
Flight Energy Savings	14.0%	15.6%	15.6%
Num. of Success Mission Increase	17.0%	18.6%	17.9%

BERRY is adaptive across environments, and consistently improves robustness and efficiency



Effectiveness Across Drone Platforms



Bitcraze Crazyflie



DJI Tello

Drone Type	Network Policy	Rotor Power	Compute Power	BERRY Flight Energy Reduction	BERRY #Missions Increase
Crazyflie	3*Conv + 2*FC	93.5%	6.5%	15.62%	18.51%
DJI Tello	3*Conv + 2*FC	97.2%	2.8%	9.91%	9.96%



Effectiveness Across Drone Platforms



Bitcraze Crazyflie



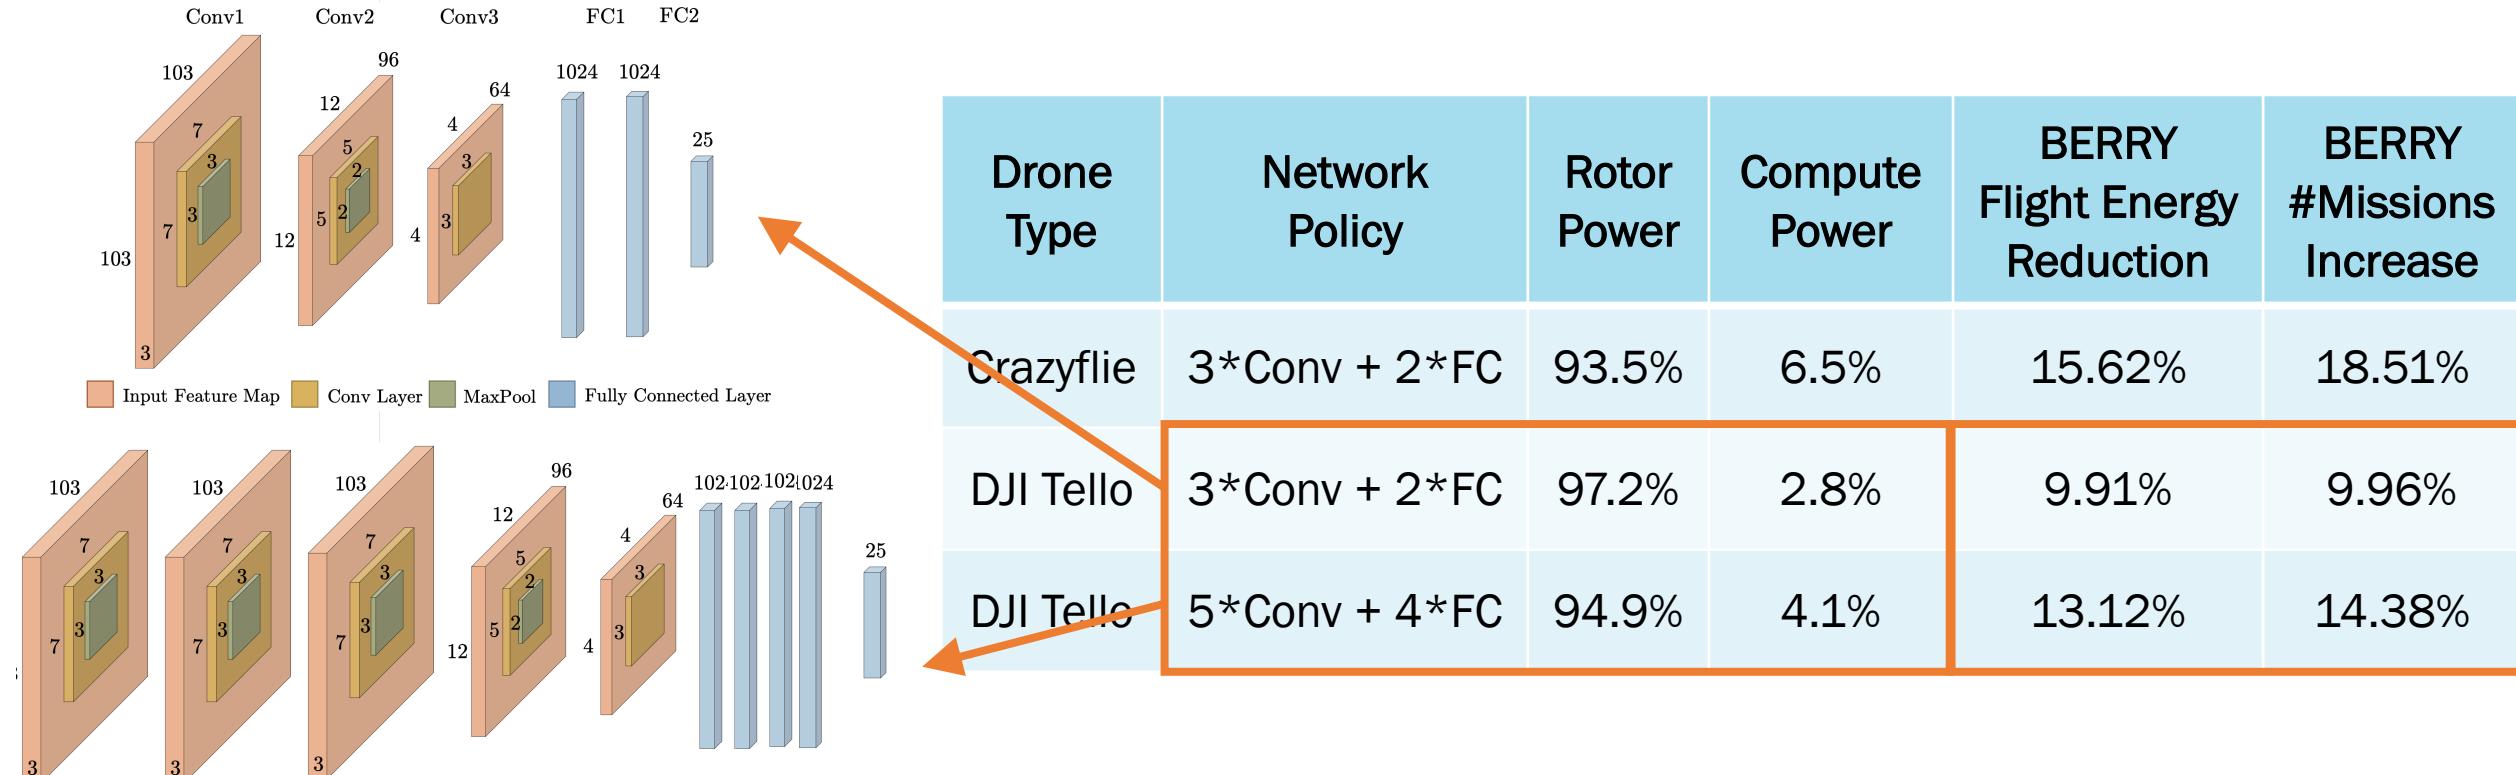
DJI Tello

Drone Type	Network Policy	Rotor Power	Compute Power	BERRY Flight Energy Reduction	BERRY #Missions Increase
Crazyflie	3*Conv + 2*FC	93.5%	6.5%	15.62%	18.51%
DJI Tello	3*Conv + 2*FC	97.2%	2.8%	9.91%	9.96%

BERRY is adaptive across drone types, and consistently improves robustness and efficiency



Effectiveness Across Model Architectures



BERRY is adaptive across models, and consistently improves robustness and efficiency



Effectiveness Across Profiled Bit Errors

Chip 1



Chip 2



Chips and Error Rates p (%)

Chip 1 (random pattern)

BERRY $p=0.5$

SuccRate SR (%)

$p=0.16$ $p=0.74$

$SR=84.0$ $SR=77.2$

Flight Energy E (J)

$p=0.16$ $p=0.74$

$E=48.46$ $E=54.63$

Chip 2 (column-aligned)

BERRY $p=0.5$

$p=0.067$ $p=0.32$

$SR=86.0$ $SR=81.8$

$p=0.067$ $p=0.32$

$E=46.98$ $E=51.27$

Baseline $p=0$ @1V

$SR=88.4$

$E=53.19$

BERRY is adaptive across bit error patterns, and consistently improves robustness and efficiency

On-Device Error-Aware Robust Learning

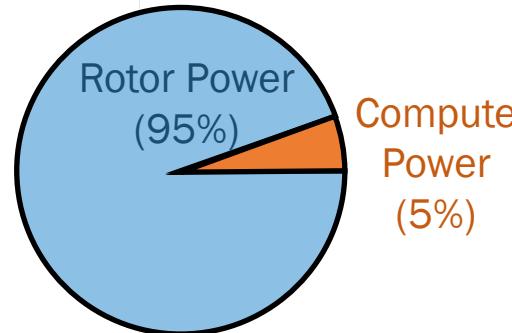
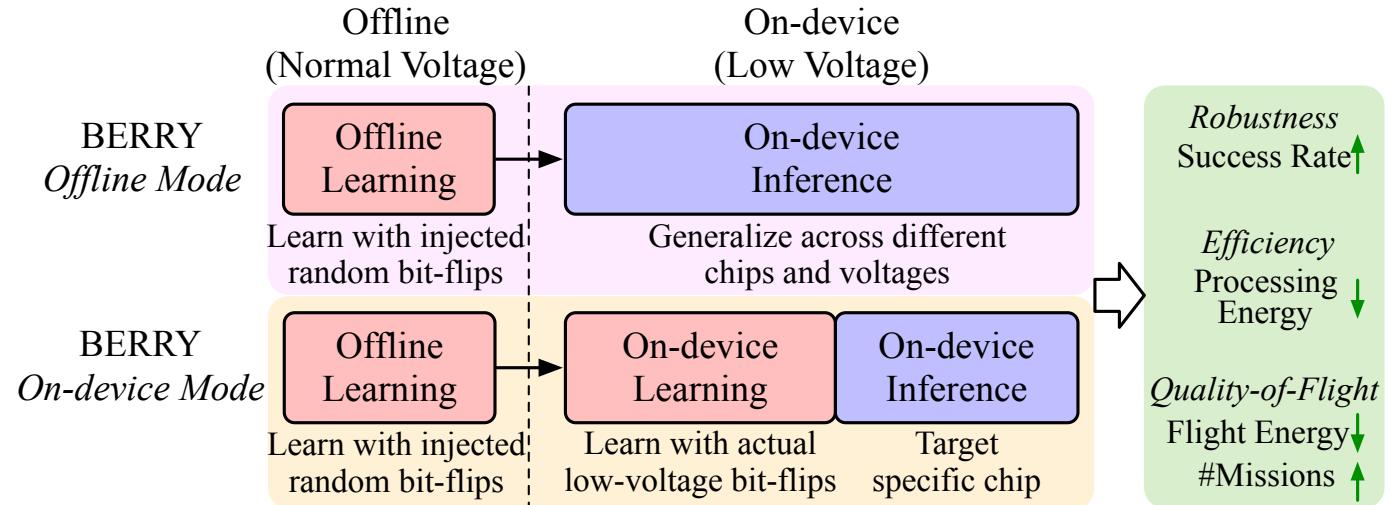
Low-Voltage Operation			Operating Efficiency	Robustness	Quality-of-Flight	
Num. of Learning Steps	Operating Voltage	Learning Energy (J)	Energy Savings	Success Rate (%)	Flight Energy (J)	Num. of Missions*
On-Device BERRY	4000	$0.77V_{min}$	1849	3.43×	84.6	264.2 48.19
		$0.70V_{min}$	1807	4.16×	82.4	266.5 46.52
	6000	$0.77V_{min}$	2775	3.43×	85.0	260.9 49.03
		$0.70V_{min}$	2711	4.16×	84.8	255.1 50.01
Offline BERRY		$0.77V_{min}$	-	3.43×	84.4	265.5 47.84
		$0.70V_{min}$	-	4.16×	63.8	375.6 25.56
Baseline	1V	-	1×	85.2	294.7	43.50

* Does not include on-device learning flight energy, evaluated for missions after learning.

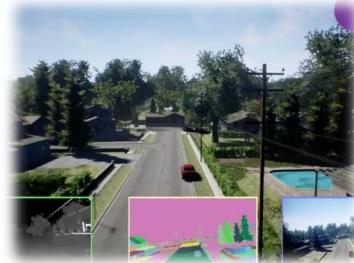
BERRY on-device robust learning consumes **on-the-fly learning energy**, but can enable **lower operating voltage, improved robustness and higher mission efficiency**.

*Please refer to paper for evaluation result details

Summary

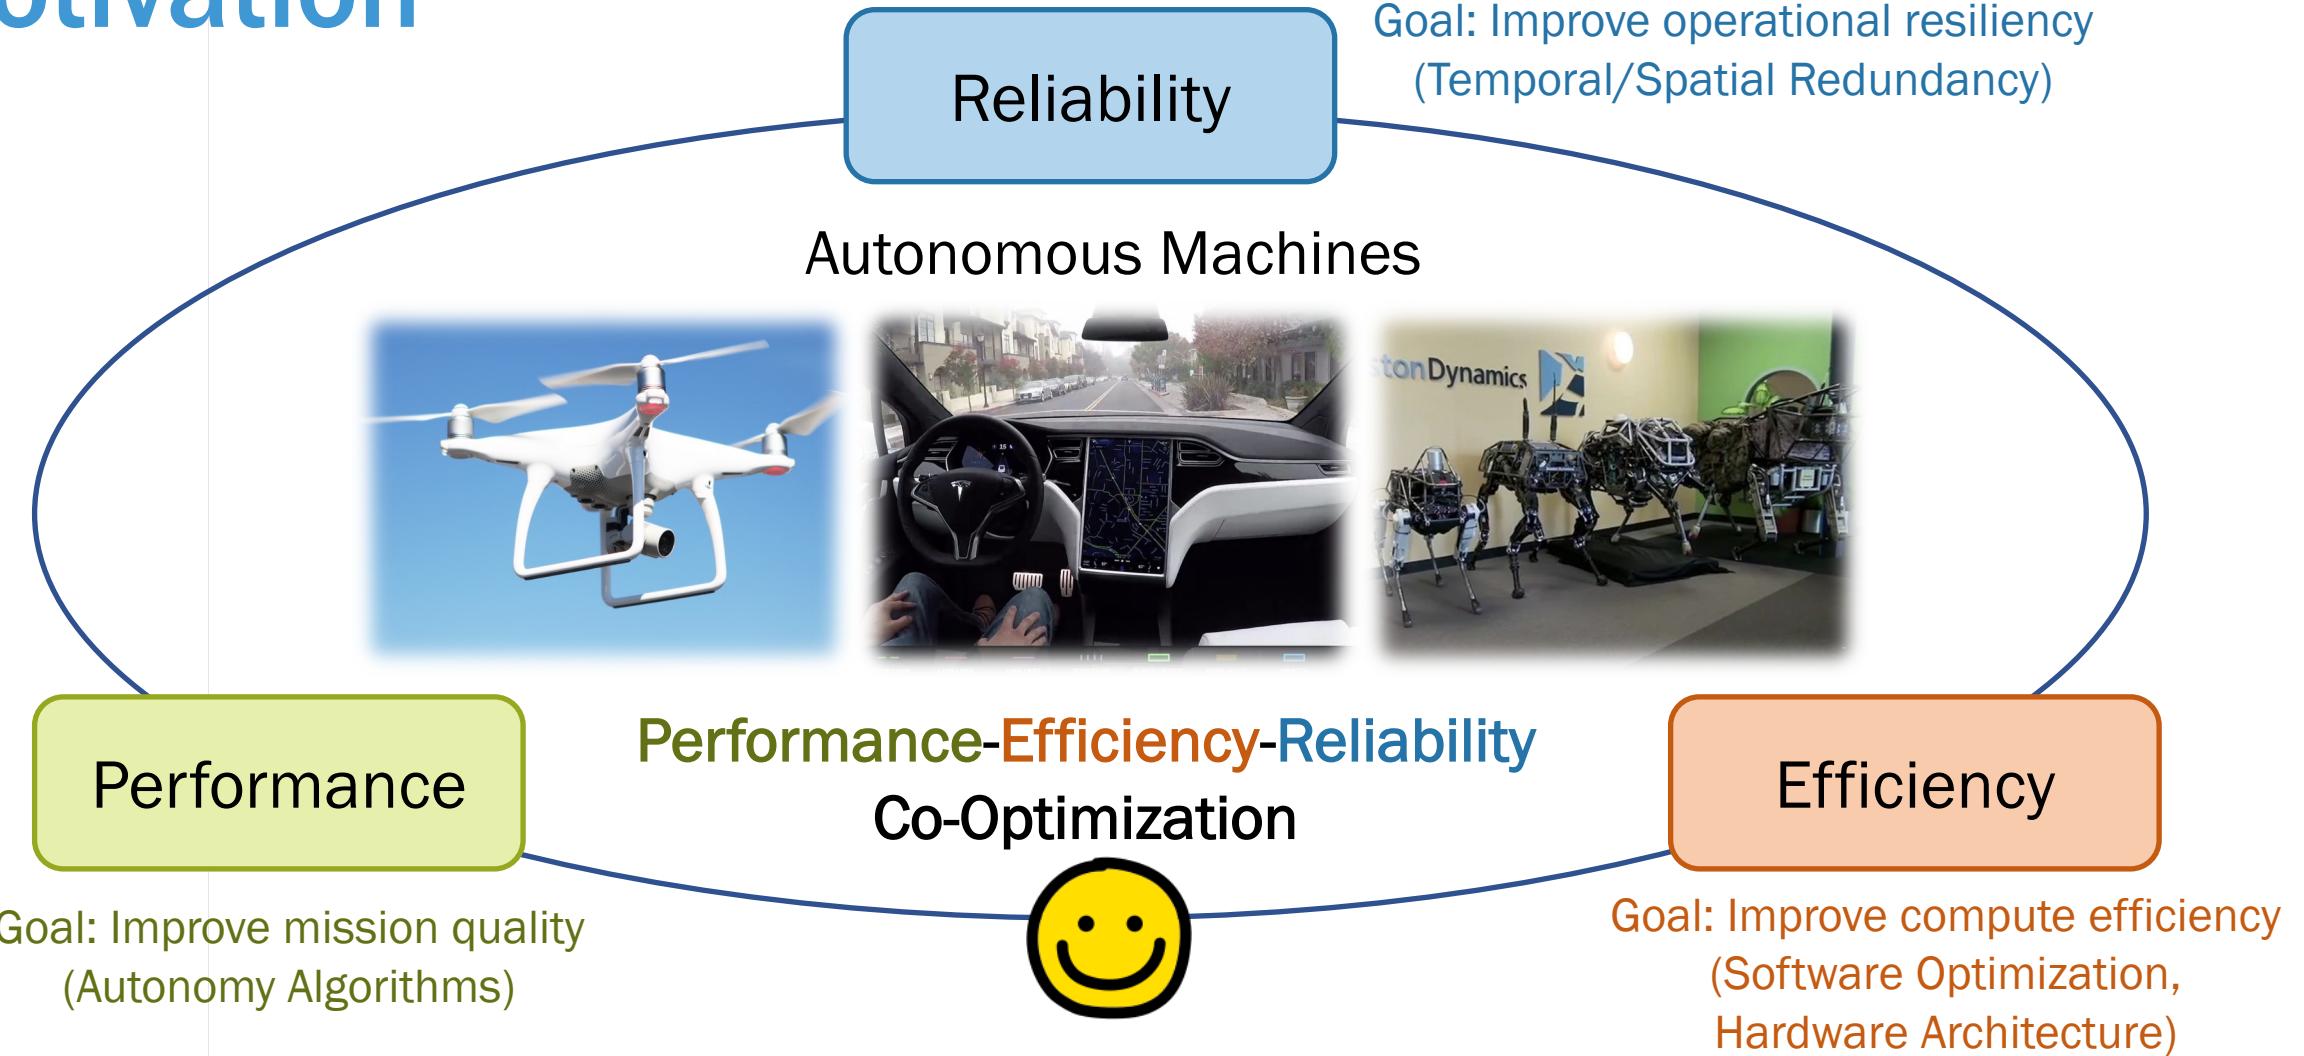


Compute power: small portion, big impact!



Aggressive energy-saving yet computational-resilient

Motivation





Semiconductor
Research
Corporation



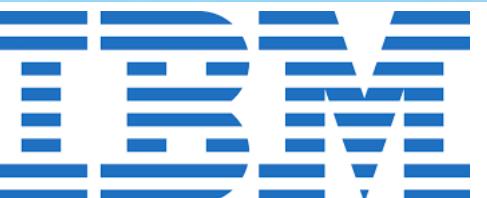
CoCoSys
CENTER FOR THE
CO-DESIGN OF COGNITIVE SYSTEMS

BERRY: Bit Error Robustness for Energy-Efficient Reinforcement Learning-Based Autonomous Systems

Zishen Wan¹, Nandhini Chandramoorthy², Karthik Swaminathan²,
Pin-Yu Chen², Vijay Janapa Reddi³, Arijit Raychowdhury¹

✉ zishenwan@gatech.edu

¹Georgia Institute of Technology ²IBM Research ³Harvard University



HARVARD
UNIVERSITY

