

NSFlow: An End-to-End FPGA Framework with Scalable Dataflow Architecture for Neuro-Symbolic AI

Hanchen Yang^{1*}, Zishen Wan^{1*}, Ritik Raj¹, Joongun Park¹, Ziwei Li¹,
Ananda Samajdar², Arijit Raychowdhury¹, Tushar Krishna¹
(*Equal Contributions)

¹Georgia Tech, Atlanta, GA, USA

²IBM Research, Yorktown Heights, NY, USA



SPONSORED BY






Background

SPONSORED BY




Background

Traditional NNs are not good at reasoning



- (i) Remove all gray spheres. How many spheres are there? (3), (ii) Take away 3 cubes. How many objects are there? (7), (iii) How many blocks must be removed to get 1 block? (2)


Complex Question Answering
NN accuracy: 50%



Abstract Reasoning
NN accuracy: 53%

IMO 2015 P3

"Let ABC be an acute triangle. Let (O) be its circumcircle, H its orthocenter, and F the foot of the altitude from A . Let M be the midpoint of BC . Let Q be the point on (O) such that $QH \perp QA$ and let K be the point on (O) such that $KH \perp KQ$. Prove that the circumcircles (O_1) and (O_2) of triangles FKM and KQH are tangent to each other."



Automated Theorem Proving
NN accuracy: 0%



Interactive Learning
NN accuracy: 71%

Scenario
Imagine that a stranger will give Hank one thousand dollars to break all the windows in his neighbor's house without his neighbor's permission. Hank carries out the stranger's request.

Imagine that there are five people who are waiting in line to use a single-occupancy bathroom at a concert venue. Someone at the back of the line needs to throw up immediately. That person skips to the front of the line instead of waiting in the back.

At a summer camp, there is a pool. Right next to the pool is a tent where the kids at the camp have art class. The camp made a rule that there would be no cannonballing in the pool so that the art wouldn't get ruined by the splashing water. Today, there is a bee attacking this kid, and she needs to jump into the water quickly. This kid cannonballs into the pool.



Ethical Decision Making
NN accuracy: 65%

Farmer John has N cows ($2 \leq N \leq 10^5$). Each cow has a breed that is either Guernsey or Holstein. As is often the case, the cows are standing in a line, numbered $1 \dots N$ in this order.

Over the course of the day, each cow writes down a list of cows. Specifically, cow i 's list contains the range of cows starting with herself (cow i) up to and including cow E_i ($i \leq E_i \leq N$).

FJ has recently discovered that each breed of cow has exactly one distinct leader. FJ does not know who the leaders are, but he knows that each leader must have a list that includes all the cows of their breed, or the other breed's leader (or both).

Help FJ count the number of pairs of cows that could be leaders. It is guaranteed that there is at least one possible pair.

Problem


Competitive Programming
NN accuracy: 8.7%



Background

Neural-Symbolic AI (NSAI)


- A **compositional system** to enhance **cognitive capability** for reasoning tasks.



Background - NSAI

Neural-Symbolic AI **example**

- **Visual Reasoning**




Question: *Are there an equal number of large things and metal spheres?*

Background - NSAI

Neural-Symbolic AI **example**

- **Visual Reasoning**



Question: Are there an **equal number** of large things and metal spheres?


3 large
things!



Background - NSAI

Neural-Symbolic AI **example**

- **Visual Reasoning**



Question: *Are there an equal number of large things and metal spheres?*

3 large things!


3 metal spheres!




Background - NSAI

Neural-Symbolic AI **example**

- **Visual Reasoning**




Question: Are there an **equal number** of large things and metal spheres?



Background - NSAI

Neural-Symbolic AI **example**


- **Visual Reasoning**



Visual Perception

Question Understanding



Question: *Are there an equal number of large things and metal spheres?*



Background - NSAI

Neural-Symbolic AI **example**


- **Visual Reasoning**



Background - NSAI

NSAI algorithm structure

Representative Neuro-Symbolic AI Workloads		Neuro-Vector-Symbolic Architecture (NVSA) [13]	Multiple-Input-Multiple-Output Neural Networks (MIMONet) [24]	Probabilistic Abduction via Learning Rules in Vector-symbolic Architecture (LVRF) [12]	Probabilistic Abduction and Execution Learner (PrAE) [40]
Compute Pattern	Neuro	CNN	CNN/Transformer	CNN	CNN
Symbolic	VSA binding/unbinding (Circular Conv)	VSA binding (Circular Conv)	VSA binding/unbinding (Circular Conv)	Probabilistic abduction	
Application Scenario	Use Case	Spatial-temporal and abstract reasoning	Multi-input simultaneously processing	Probabilistic reasoning, OOD data processing	Spatial-temporal and abstract reasoning
Advantage vs. Neural	Higher joint representation efficiency, Better reasoning capability, Transparency	Higher throughput, Lower latency, Compositional compute, Transparency	Stronger OOD handling capability, One-pass learning, Higher flexibility, Transparency	Higher generalization, Transparency, Interpretability, Robustness	



Motivation




SPONSORED BY



Motivation


NSAI workloads system characterization



Z. Wan et al., "Towards Cognitive AI Systems: Workload and Characterization of Neuro-Symbolic AI,"
2024 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)




Motivation

NSAI workloads system characterization






Motivation

NSAI workloads system characterization



Motivation

NSAI workloads system characterization



Neuro-symbolic workload exhibits high latency compared to neural models;
Symbolic component is processed inefficiently on off-the-shelf CPU/GPUs

Motivation

NSAI algorithms deployment challenges:

- **Memory & compute inefficiency**



Motivation

NSAI algorithms deployment challenges:

- **Memory & compute inefficiency**
- **Heterogeneous compute kernel**



Motivation

NSAI algorithms deployment challenges:


- **Memory & compute inefficiency**
- **Heterogeneous compute kernel**
- **Algorithm diversity**




Motivation

NSAI algorithms deployment challenges:

- **Memory & compute inefficiency**
- **Heterogeneous compute kernel**
- **Algorithm diversity**



Efficient and flexible architecture



Motivation


NSAI algorithms deployment challenges:

- **Memory & compute inefficiency**
- **Heterogeneous compute kernel**
- **Algorithm diversity**




Efficient and flexible architecture

Automated and customized
FPGA deployment





Motivation

NSFlow



Motivation

NSFlow



NSFlow Framework




SPONSORED BY



NSFlow Framework


An **end-to-end** automated **FPGA** framework for accelerating and deploying **generic NSAI** workloads.



NSFlow Framework

An **end-to-end** automated **FPGA** framework for accelerating and deploying **generic NSAI** workloads.


- **Identifies data dependency** for the workload



NSFlow Framework

An **end-to-end** automated **FPGA** framework for accelerating and deploying **generic NSAI** workloads.


- **Identifies data dependency** for the workload
- **Explores design space** with parameterizable HW blocks.



NSFlow Framework

An **end-to-end** automated **FPGA** framework for accelerating and deploying **generic NSAI** workloads.


- **Identifies data dependency** for the workload
- **Explores design space** with parameterizable HW blocks.
- Generates and deploys **optimal dataflow architecture** on FPGA



NSFlow Framework


An **end-to-end** automated **FPGA** framework for accelerating and deploying **generic NSAI** workloads.

- **Identifies data dependency** for the workload
 - **Explores design space** with parameterizable HW blocks.
 - Generates and deploys **optimal dataflow architecture** on FPGA
- ✓ *Enables automated, efficient and scalable dataflow and architecture solutions for NSAI*




NSFlow Framework

Frontend




NSFlow Framework

Frontend



NSFlow Framework

Frontend




Backend

NSFlow Framework


Frontend

Backend




NSFlow Framework

Frontend



NSFlow Framework

Frontend



NSFlow Backend

Flexible Hardware Architecture




SPONSORED BY



NSFlow Backend

An accelerator template

w/ RTL blocks **parameterizable** by config files




NSFlow Backend

An accelerator template


w/ RTL blocks **parameterizable** by config files

- ❑ **Adaptive Systolic Array:**

for efficient **Neuro** & **Symbolic** processing




NSFlow Backend





$$\text{Example: } (A_1, A_2, A_3) \odot (B_1, B_2, B_3) = (A_1B_1 + A_2B_2 + A_3B_3, A_1B_3 + A_2B_1 + A_3B_2, A_1B_2 + A_2B_3 + A_3B_1)$$

■ Stationary Reg. ■ Passing Reg ■ Streaming Reg. ■ Partial Sum Reg. ○ MAX



Z. Wan, H. Yang, R. Raj et al., "CogSys: Efficient and Scalable Neurosymbolic Cognition System via Algorithm-Hardware Co-Design," 2025 IEEE International Symposium on High Performance Computer Architecture (HPCA)

NSFlow Backend



Enables efficient **Vector Symbolic** operation processing on systolic array

NSFlow Backend

An accelerator template


w/ RTL blocks **parameterizable** by config files

- ❑ **Adaptive Systolic Array:**

for efficient **Neuro** & **Symbolic** processing

- ❑ **SIMD Unit:**

for **element-wise ops**



NSFlow Backend

An accelerator template

w/ RTL blocks **parameterizable** by config files

- ❑ **Adaptive Systolic Array:**


for efficient **Neuro** & **Symbolic** processing

- ❑ **SIMD Unit:**

for **element-wise ops**

- ❑ **BRAM Blocks:**

for flexible on-chip memory



NSFlow Backend

An accelerator template

w/ RTL blocks **parameterizable** by config files

- ❑ **Adaptive Systolic Array:**

for efficient **Neuro** & **Symbolic** processing

- ❑ **SIMD Unit:**


for **element-wise ops**

- ❑ **BRAM Blocks:**

for flexible on-chip memory

- ❑ **Control Logic:**

for HW-level task scheduling



NSFlow Frontend

Dataflow Architecture Generation (DAG)



SPONSORED BY



NSFlow Frontend

- Analyze workload **characteristics** and **data dependencies**
- Explore optimal HW **configurations** and **dataflow** for backend



NSFlow Frontend

1. Extract **Execution Trace** from the workload:

wokload.py



NSFlow Frontend

1. Extract Execution Trace from the workload:


wokload.py => et.json

```
graph():
    ...
    // Neuro Operation - CNN (Resnet18)
    %relu_1[16, 64, 160, 160] : call_module[relu](args = (%bn1
        [16, 64, 160, 160]))
    %maxpool_1[16, 64, 160, 160] : call_module[maxpool](args =
        (%relu_1[16, 64, 160, 160]))
    %conv2d_1[16, 64, 160, 160] : call_module[conv2d](args =
        (%maxpool_1[16, 64, 160, 160]))
    ...
    // Symbolic Operations
    // Inverse binding of two block codes vectors by
        blockwise circular correlation
    %inv_binding_circular_1[1, 4, 256] : call_function[nvs.
        inv_binding_circular](args = (%vec_0[1, 4, 256], %
        vec_1[1, 4, 256]))
    %inv_binding_circular_2[1, 4, 256] : call_function[nvs.
        inv_binding_circular](args = (%vec_3[1, 4, 256], %
        vec_4[1, 4, 256]))
    // Compute similarity between two block codes vectors
    %match_prob_1[1] : call_function[nvs.match_prob](args
        = (%inv_binding_circular_1[1, 4, 256], %vec_2
        [1, 4, 256]))
    // Compute similarity between a dictionary and a batch
        of query vectors
    %match_prob_multi_batched_1[1]: call_function[nvs.
        match_prob_multi_batched](args = (%
        inv_binding_circular_2[1, 4, 256], %vec_5[7, 4, 256]))
    %sum_1[1] : call_function[torch.sum](args = (%
        match_prob_multi_batched_1[1]))
    %clamp_1[1] : call_function[torch.clamp](args = (%sum_1
        [1]))
    %mul_1[1] : call_function[operator.mul](args = (%
        match_prob_1[1], %clamp_1[1]))
    ...
```



NSFlow Frontend


2. Generate **Dataflow Graph** for deployment:



1 Perform DFS in the execution graph, and identify **critical path** for a single run.


NSFlow Frontend

2. Generate **Dataflow Graph** for deployment:




NSFlow Frontend

2. Generate Dataflow Graph for deployment:



NSFlow Frontend

2. Generate Dataflow Graph for deployment:



NSFlow Frontend

3. Explore optimal **HW config** and **array partition** strategy



NSFlow Frontend

3. Explore optimal HW config and array partition strategy

Algorithm 1: NSFlow Two-Phase DSE Algorithm

Data: $R_l, R_v, Range_H$ (H search range), $Range_W$ (W search range), M (max #PEs), $Iter_{max}$ (Phase II max iterations)
Result: H, W, N (total #sub-arrays), N_l, N_v

```
1 /* Phase I */
2 for  $H$  in  $Range_H$ ,  $W$  in  $Range_W$  do
3    $N = \lfloor M / (H \times W) \rfloor$  // get total #sub-arrays
4   for  $\bar{N}_l$  in  $[1, N]$  do
5     // get optimal HW config for parallel mapping
6     Set all elements in  $N_l$  to  $\bar{N}_l$ 
7     Set all elements in  $N_v$  to  $N - \bar{N}_l$ 
8      $t_{para} = \max(t_{nn}(H, W, N_l), t_{vsa}(H, W, N_v))$ 
9     Save the  $H, W, \bar{N}_l$  (and  $\bar{N}_v$ ) with minimal  $t_{para}$ .
10    end
11    // get sequential runtime
12     $t_{seq} = \sum_i^G f_{l_i}(H, W, N) +$ 
13       $\min(\sum_j^G f_{v_j, temp}(H, W, N), \sum_j^G f_{v_j, spatial}(H, W, N))$ 
14    // Set to sequential mode in case it has better performance
15    Return and set sequential mode if  $t_{seq} < t_{para}$  else Continue
16  end
17 /* Phase II */
18 for  $it$  in  $Iter_{max}$  do
19   for layer  $i$  in  $R_l$  do
20     Locate VSA node  $j'$  and  $j''$  where layer  $i$  starts and ends
21     if  $t_{seq} < t_{para}$  do  $N_l[i] --$ ;  $N_v[j' : j''] ++$ ;
22     else do  $N_l[i] ++$ ;  $N_v[j' : j''] --$ ;
23      $t_{para} = \max(t_{nn}(H, W, N_l), t_{vsa}(H, W, N_v))$ 
24     Save the  $H, W, N_l, N_v$  with minimal  $t_{para}$ .
25   end
26 end
27 Return  $H, W, N, N_l, N_v$ .
```



NSFlow Frontend

3. Explore optimal HW config and array partition strategy

- *Phase I:* Assuming static partition, find the optimal array size (H, W, N).

Algorithm 1: NSFlow Two-Phase DSE Algorithm

Data: R_l , R_v , Range_H (H search range), Range_W (W search range), M (max #PEs), Iter_{\max} (Phase II max iterations)
Result: H , W , N (total #sub-arrays), N_l , N_v

```
1 /* Phase I */
2 for  $H$  in  $\text{Range}_H$ ,  $W$  in  $\text{Range}_W$  do
3    $N = \lfloor M / (H \times W) \rfloor$  // get total #sub-arrays
4   for  $\bar{N}_l$  in  $[1, N]$  do
5     // get optimal HW config for parallel mapping
6     Set all elements in  $N_l$  to  $\bar{N}_l$ 
7     Set all elements in  $N_v$  to  $N - \bar{N}_l$ 
8      $t_{para} = \max(t_{nn}(H, W, N_l), t_{vsa}(H, W, N_v))$ 
9     Save the  $H, W, \bar{N}_l$  (and  $\bar{N}_v$ ) with minimal  $t_{para}$ .
10    end
11    // get sequential runtime
12     $t_{seq} = \sum_i^G f_{l_i}(H, W, N) +$ 
13       $\min(\sum_j^G f_{v_j, temp}(H, W, N), \sum_j^G f_{v_j, spatial}(H, W, N))$ 
14    // Set to sequential mode in case it has better performance
15    Return and set sequential mode if  $t_{seq} < t_{para}$  else Continue
16  end
17  /* Phase II */
18  for  $it$  in  $\text{Iter}_{\max}$  do
19    for layer  $i$  in  $R_l$  do
20      Locate VSA node  $j'$  and  $j''$  where layer  $i$  starts and ends
21      if  $t_{seq} < t_{para}$  do  $N_l[i] --$ ;  $N_v[j' : j''] ++$ ;
22      else do  $N_l[i] ++$ ;  $N_v[j' : j''] --$ ;
23       $t_{para} = \max(t_{nn}(H, W, N_l), t_{vsa}(H, W, N_v))$ 
24      Save the  $H, W, N_l, N_v$  with minimal  $t_{para}$ .
25    end
26  end
27 Return  $H, W, N, N_l, N_v$ .
```



NSFlow Frontend

3. Explore optimal HW config and array partition strategy

- *Phase I:* Assuming **static partition**, find the optimal array size (H, W, N).
- *Phase II:* Fine-tune for **dynamic partition** to better balance **Neuro** & **Symb** at runtime

Algorithm 1: NSFlow Two-Phase DSE Algorithm

Data: R_l , R_v , Range_H (H search range), Range_W (W search range), M (max #PEs), Iter_{\max} (Phase II max iterations)
Result: H , W , N (total #sub-arrays), N_l , N_v

```
1 /* Phase I */
2 for  $H$  in  $\text{Range}_H$ ,  $W$  in  $\text{Range}_W$  do
3    $N = \lfloor M / (H \times W) \rfloor$  // get total #sub-arrays
4   for  $\bar{N}_l$  in  $[1, N]$  do
5     // get optimal HW config for parallel mapping
6     Set all elements in  $N_l$  to  $\bar{N}_l$ 
7     Set all elements in  $N_v$  to  $N - \bar{N}_l$ 
8      $t_{para} = \max(t_{nn}(H, W, N_l), t_{vsa}(H, W, N_v))$ 
9     Save the  $H, W, N_l$  (and  $\bar{N}_v$ ) with minimal  $t_{para}$ .
10 end
11 // get sequential runtime
12  $t_{seq} = \sum_i^G f_{l_i}(H, W, N) +$ 
13    $\min(\sum_j^G f_{v_j, temp}(H, W, N), \sum_j^G f_{v_j, spatial}(H, W, N))$ 
14 // Set to sequential mode in case it has better performance
15 Return and set sequential mode if  $t_{seq} < t_{para}$  else Continue
16 /* Phase II */
17 for  $it$  in  $\text{Iter}_{\max}$  do
18   for layer  $i$  in  $R_l$  do
19     Locate VSA node  $j'$  and  $j''$  where layer  $i$  starts and ends
20     if  $t_{seq} < t_{para}$  do  $N_l[i] --$ ;  $N_v[j' : j''] ++$ ;
21     else do  $N_l[i] ++$ ;  $N_v[j' : j''] --$ ;
22      $t_{para} = \max(t_{nn}(H, W, N_l), t_{vsa}(H, W, N_v))$ 
23     Save the  $H, W, N_l, N_v$  with minimal  $t_{para}$ .
24   end
25 end
26 Return  $H, W, N, N_l, N_v$ .
```



NSFlow Frontend

3. Explore optimal HW config and array partition strategy

- *Phase I:* Assuming **static partition**, find the optimal array size (H, W, N).
- *Phase II:* Fine-tune for **dynamic partition** to better balance **Neuro** & **Symb** at runtime

✓ Reduces search space $\times 10^{100}$

	HW config (H, W, N)	Array partition and mapping	Total design space, $m = 10$
Original	$m \times (m + 1)/2$	$(N - 1)^k$ for each N	10^{300}
DAG	$PhaseI: 1/4 \leq H/W \leq 16$	$PhaseII: \text{Iter} \times \#layers$	10^3

Algorithm 1: NSFlow Two-Phase DSE Algorithm

Data: $R_l, R_v, Range_H$ (H search range), $Range_W$ (W search range), M (max #PEs), $Iter_{max}$ (Phase II max iterations)
Result: H, W, N (total #sub-arrays), N_l, N_v

```
1 /* Phase I */
2 for  $H$  in  $Range_H$ ,  $W$  in  $Range_W$  do
3    $N = \lfloor M/(H \times W) \rfloor$  // get total #sub-arrays
4   for  $\bar{N}_l$  in  $[1, N]$  do
5     // get optimal HW config for parallel mapping
6     Set all elements in  $N_l$  to  $\bar{N}_l$ 
7     Set all elements in  $N_v$  to  $N - \bar{N}_l$ 
8      $t_{para} = \max(t_{nn}(H, W, N_l), t_{vsa}(H, W, N_v))$ 
9     Save the  $H, W, \bar{N}_l$  (and  $\bar{N}_v$ ) with minimal  $t_{para}$ .
10 end
11 // get sequential runtime
12  $t_{seq} = \sum_i^G f_{l_i}(H, W, N) +$ 
13    $\min(\sum_j^G f_{v_j, temp}(H, W, N), \sum_j^G f_{v_j, spatial}(H, W, N))$ 
14 // Set to sequential mode in case it has better performance
15 Return and set sequential mode if  $t_{seq} < t_{para}$  else Continue
16 /* Phase II */
17 for  $it$  in  $Iter_{max}$  do
18   for layer  $i$  in  $R_l$  do
19     Locate VSA node  $j'$  and  $j''$  where layer  $i$  starts and ends
20     if  $t_{seq} < t_{para}$  do  $N_l[i] --$ ;  $N_v[j' : j''] ++$ ;
21     else do  $N_l[i] ++$ ;  $N_v[j' : j''] --$ ;
22      $t_{para} = \max(t_{nn}(H, W, N_l), t_{vsa}(H, W, N_v))$ 
23     Save the  $H, W, N_l, N_v$  with minimal  $t_{para}$ .
24   end
25 end
26 Return  $H, W, N, N_l, N_v$ .
```



Evaluation



SPONSORED BY



Evaluation

Experiments setup


- Workloads:
 - Algorithms: *NVSA*, *MIMONet*, *LVRF*
 - Datasets: *RAVEN*, *I-RAVEN*, *PGM*, *CVR*, and *SVRT*
- Hardwares:
 - Baselines: *TX2*, *Xavier NX*, *Xeon CPU*, *RTX 3080*, *ML accelerators (TPU, Xilinx DPU)*
 - FPGA deployment: *AMD U250*

Workloads	Precision		AdArray Configuration		SIMD Size	On-chip SRAM Blocks (BRAM)			On-chip Cache (URAM)	AMD U250 Utilization						Frequency
	NN	Symb	Size (H, W, N)	Default Partition ($\bar{N}_l : \bar{N}_v$)		Mem A1, Mem A2	Mem B	Mem C		DSP	LUT	FF	BRAM	URAM	LUTRAM	
NVSA	INT8	INT4	32, 16, 16	14 : 2	64	2.7 MB, 1.1 MB	2.7 MB	1.6 MB	16.2 MB	89%	56%	60%	34%	8%	24%	272 MHz
MIMONet	INT8	INT8	32, 32, 8	6 : 2	64	3.4 MB, 1.2 MB	3.4 MB	2.1 MB	20.1 MB	89%	44%	52%	43%	10%	20%	272 MHz
LVRF	INT8	INT4	32, 16, 16	14 : 2	64	2.7 MB, 0.96 MB	2.7 MB	1.4 MB	15.5 MB	89%	56%	60%	31%	7%	24%	272 MHz

Evaluation

End-to-end runtime improvement

✓ ~2x speedup over GPU, 2~8x speedup over TPU, ~3x speedup over DPU



Evaluation

Mixed-precision optimization


✓~6x Memory reduction

Reasoning Accuracy	FP32	FP16	INT8	MP (IN8 for NN, INT4 for Symb)	INT4
RAVEN [39]	98.9%	98.9%	98.7%	98.0%	92.5%
I-RAVEN [16]	99.0%	98.9%	98.8%	98.1%	91.3%
PGM [3]	68.7%	68.6%	68.4%	67.4%	59.9%
Memory	32MB	16MB	8MB	5.5MB	4MB

Evaluation

Scalability

- ✓ Only 4x runtime increase when symbolic workloads scale by 150x



Conclusion



SPONSORED BY



Conclusion

NSFlow is the first **end-to-end design automation** framework dedicated to accelerate **generic NSAI** systems

Conclusion

NSFlow is the first **end-to-end design automation** framework dedicated to accelerate **generic NSAI** systems

- *Identifies the unique optimization opportunities for NSAI acceleration*

Conclusion

NSFlow is the first **end-to-end design automation** framework dedicated to accelerate **generic NSAI** systems

- *Identifies the unique optimization opportunities for NSAI acceleration*
- *Explores the dataflow and architecture design space with a novel algorithm*

Conclusion

NSFlow is the first **end-to-end design automation** framework dedicated to accelerate **generic NSAI** systems

- *Identifies the unique optimization opportunities for NSAI acceleration*
- *Explores the dataflow and architecture design space with a novel algorithm*
- *Generates a efficient scalable design for FPGA deployment*

Conclusion

NSFlow paves the way
for advancing efficient cognitive reasoning systems and
unlocking new possibilities in NSAI.

Thank you!



SPONSORED BY

